


COMP11212 Fundamentals of Computation

Part 1: Formal Languages

Sean Bechhofer
School of Computer Science
University of Manchester
sean.bechhofer@manchester.ac.uk

2020-01-20



Organisational issues

When, what, where

This course will be taught as follows.

- Lectures: Will take place on Tuesdays and Wednesdays at 11:00 in Chemistry G.51.¹.
- Examples classes: Will take place from Week 2 as follows.

Groups	Time	Location
M+W	Tuesday 09.00	Collab 1
X	Wednesday 10.00	Collab 1
Y	Monday 10.00	Collab 1
Z	Thursday 09.00	G41

Two parts

This course consists of two distinct parts, taught by two members of staff, Sean Bechhofer and Gareth Henshall.

Lectures	Week	Content	Taught by
1	1	Intro + start Part 1	Sean & Gareth
2–10	1–5	Part 1	Sean
11–20	6–10	Part 2	Gareth
21–22	11	Revision	Sean & Gareth

Examples classes in Weeks	belongs to
2–6	Part 1
7–12	Part 2

Assessment

There is *assessed coursework* in the form of exercises you are expected to complete. The coursework is worth *20% of the overall mark for the course unit*.

The exam will be a *hybrid* exam with questions set online along with questions to be answered on paper. Students have to answer **all** questions. The exam will be worth *80% of your mark for this course unit*.

¹Note that these times/locations are correct as of 2020-01-20. Please check the School timetable for confirmation in case of last minute timetable changes.

Coursework

As with other course units, the exercises use a mix of *formative* and *summative* assessment. There are exercises for you to do every week. In Weeks 3, 6, 9 and 11 there will also be some *summative* assessments and the marks awarded there will contribute towards your final coursework mark.

The weekly formative assessments will include questions to be completed online via Blackboard and paper exercises with feedback in the examples classes. The formative assessments will help to complete the summative assessments.

Online Exercises

Each week there will be online exercises to be completed in Blackboard. This will (usually) consist of three questions broken down into sub parts. Questions will be a mixture of true/false, multiple choice, multiple answer (where there may be more than one correct answer), fill in the blank and matching questions.

You should aim to complete the work for the exercises on your own. The problems are intended to help with your understanding of the material. If you do not understand that material, you will struggle to pass the exam at the end of the unit. Experience shows that students who engaged with coursework during the year performed well in the final exams.

You will receive feedback on formative paper exercises in the examples classes. Summative paper exercises should be submitted via blackboard. Solutions and feedback will be made available after the submission deadlines. You can then discuss any issues or problems you may have had in the following examples class.

The online exercises will be of various forms:

Multiple Choice (MCQ) For multiple choice questions, there will be several options. Only one option is correct, and you must select the correct option to obtain a mark.

Multiple Answer (MAQ) For multiple answer questions, there will be several options. At least one option will be correct and at least one option will be incorrect, but there *may* be more than one answer that is correct. For an MAQ, you will receive credit for selecting a correct answer and a penalty for selecting an incorrect answer.

Matching Questions (MQ) For matching questions, there will be two lists of items. You must match up the items in each list. For an MQ, *all* items must be matched correctly to obtain the marks.

Assertion-Reason (AR) Assertion-Reason questions are a particular form of MCQ. You will be given a statements:

Assertion BECAUSE Reason.

The question will ask about the correctness of the individual statements (the assertion and reason), and whether the reasoning claimed is valid. For example:

Sean teaches COMP11212 *BECAUSE* Sean likes drinking whisky.

The question will provide several alternatives:

- A The assertion and reason are both correct and the reason is valid.
- B The assertion and reason are both correct but the reason is invalid.
- C The assertion is correct, but the reason is incorrect.
- D The assertion is incorrect, but the reason is correct.
- E Both the answer and the reason are incorrect.

In this case, the answer is B, both statements are true, but the first does not follow from the second. Sean’s teaching duties are unconnected with his love of single malts.

The online exercises are formative, and the marks do not contribute towards your final grade for the unit. You are strongly recommended to do the exercises though. They give you an opportunity to practice for the kinds of question that you will face in the online sections of the final exam and will also help you judge your current understanding of the materials. There are no specific deadlines for the online exercises, so you can do them when you want, and return and re-do them as revision. We strongly recommend that you try and do them in the weeks that they are set though, as this will then help to reinforce the material being taught.

Examples Classes

Sheets will be published with exercises to be completed for each examples class. You are expected to prepare exercises **before the examples class**. During the examples class the paper exercises will be marked, you will get feedback on what you did, and you will also get help with any questions you could not do.

Marking

Formative

You will be given a “mark” in the examples class sessions to help give you an idea of your level of understanding. These marks will appear in SPOT, the department’s system that allows you to track your coursework. The numbers we will use have the following meaning:

- 0:** You attended the examples class and interacted with a GTA, but you have not tried to solve any exercise in preparation.
- 1:** You have tried to solve some exercises, but you either did not do much work or your solutions were significantly lacking in quality or accuracy.
- 2:** You have tried to solve (at least) all the core exercises for the week and your solutions are mostly solid.
- 3:** You have tried to solve all the core exercises and had a go at the stretch ones; your solutions for the core exercises are solid, and you can fix most of your mistakes for those with a little bit of prompting from the GTA.
- 4:** You tried to solve all of the exercises and your solutions are correct with only minor shortcomings.

Summative

For each of the four summative exercises, there will be four sub-exercises. Each exercise is **marked out of two**, where the marks mean the following.

- 0:** You did not make a serious attempt to solve the question.
- 1:** You made a serious attempt but did not get close to finishing the question.
- 2:** You worked your way through the exercise and made no serious mistakes.

Each summative sheet will thus be worth up to 8 marks.

The rules for marking are as follows:

- To get a mark for any exercise your solution should demonstrate that you have an understanding of the material.

- It is possible to get one mark for an exercise where you haven't made much, or any, progress, provided your rough works give evidence that you have seriously tried it. We there to be at least a page or more of your trying things out to find an answer. If you submit just a final answer for an exercise that requires steps in between you will not get the marks for it.

Solutions to the assessed exercises will be released shortly after the deadline. Because of this, there is **no** opportunity for late submission.

Changes

Note that, as with all courses, the unit has evolved over the last few years. You should thus be careful about taking advice from students who have taken the unit in previous years, as their experience and information may no longer be relevant or valid.

The key change this year is a reduction in the number of summative exercises. This does not mean that we expect you do to less work! The fact that the formative exercises do not contribute to your unit mark means that you can use them as a realistic opportunity to test your own understanding without penalty. If you find that you are not able to complete those exercises, then you may need to be spending more time trying to understand the material.

About Part 1 of the course

As detailed on the previous page, Part 1 of this course unit is concerned with *formal languages* and is taught by Sean Bechhofer.

Intended Learning Outcomes

Following completion of Part 1 of this unit, a student will be able to:

- describe formal languages using a variety of mechanisms;
- define classes of languages and demonstrate translations between those classes; and
- state key properties of classes of languages and determine when those properties hold.

These ILOs will be assessed through both coursework and exam. Note that there are further ILOs associated with Part II of the unit. See the unit web pages for details.

The following concerns organisational issues just for this first part.

Learning the material

The most important aims of this part of the course are to teach you

- a few concepts (such as formal language, regular expression, finite state automaton, context-free grammar); and
- how to solve a number of typical problems,

with an emphasis on the latter.

Exercises. The best way to learn how to solve problems is to do it, which is why the notes contain a great number of exercises. Those not covered in the exercise sheets provide further opportunities to practise, for example when revising for the exams. The exercises are part of the examinable material, although I won't ask you to recall any particular exercise in the exam. Some exercises have a *, which means that they are not part of the examinable material. Note that they are not necessarily difficult, although some of them are.

If you find you have completed all of the exercises and want more examples to test your knowledge on, try making some up with your colleagues. For example, take a DFA and convert it to a regular expression. Give it to a colleague and get them to convert it back to a regular expression. Do you get the original DFA? Would you expect to? If you don't is it equivalent? Considering these questions should help you gain an understanding of the material. The recommended textbooks also have some exercises although not all of these have published solutions.

Lectures. I use the lectures to introduce the main ideas, and to explain the important algorithms, as well as suggesting some ways for solving typical problems.

Self-study. The course has only three contact hours a week, and you are expected to spend a substantial amount of time each week working on the material by studying the notes yourself, connecting the written material with what was introduced in the lectures, and solving exercises.

Notes. The notes are fairly substantial because they aim to cover the material in detail. Typically, for each new idea or algorithm, there is a formal description as well as one or two examples. For

most people it is worthwhile trying to understand the example before tackling the general case, although in the notes the general case is sometimes described before the example (because it can be difficult to explain the idea using just a special case). Note that there is a glossary that you can use to quickly find the meaning of the various technical terms used.

Mathematical notation. The language used to describe much of theoretical computer science is that of mathematics, because that is the *only* way of describing with precision what we are talking about. In the main part of the notes I have tried to remain as informal as possible in my descriptions, although all the definitions given are formal. You will not be required to repeat these definitions in the exam, you only need to understand informally what they mean. There is an appendix which spells out in more detail how the language of mathematics is used to make rigorous the concepts and operations we use. If you only understand the informal side of the course you can pass the exam (and even get a very good mark) but to master the material (and get an excellent result in the exam) you will have to get to grips with the mathematical notation as well. A maximum of 10% of the final mark depends on the material in the Appendix.

Examples classes. The examples classes give you the opportunity to get help with any problems you have, either with solving exercises or with understanding the notes. There are five examples classes associated with this part of the course, in Weeks 2–6. For each of these you are expected to prepare by solving the paper-based exercises and completing online exercises. **Solutions** for each exercise sheet will be made available online after the last associated examples class has taken place².

Marking criteria. The marking criteria for the assessed coursework are *stricter* than those for the exam. In particular, I ask you to follow various algorithms as described in the notes, whereas in the exam I'm happy if you can solve the problem at hand, and I don't mind how exactly you do that. In all cases it is important to show your work—if you only give an answer you may lose a substantial number of marks.

Revision. For revision purposes I suggest going over all the exercises again. Some of the exercises will probably be new to you since they are not part of the set preparation work for the examples classes. Also, you can turn all the NFAs you encounter along the way into DFAs (removing ϵ -transitions as required). Finally, there are exams from previous years to practice on, although in keeping with School policy, the solutions to the exams are not published. If you can do the exercises you will do well in the exam. You will not be asked to repeat definitions, but knowing about properties of regular and context-free languages may be advantageous. You should also be aware of the few theorems, although you will not be asked to recite them.

Webpage The course has a web page, where notes, information and relevant links can be found: <http://studentnet.cs.manchester.ac.uk/ugt/COMP11212/>

Reading

For this part of the course these notes cover all the examinable material. However, sometimes it can be useful to have an additional source to see the same material introduced and explained in a slightly different way. Also, there are new examples and exercises to be looked at. For this purpose you may find one of the following useful.

M. Sipser. **Introduction to the Theory of Computation.** *PWS Publishing Company*, 1997. ISBN 0-534-94728-X.

This is a fairly mathematical book that nonetheless tries to keep mathematical notation at a minimum. It contains many illustrated examples and aims to explain ideas rather than going through proofs mechanically. A 2005 edition is also available. Around £50—this book is very well thought of and even used copies are quite expensive. Relevant to this course: Chapters 0 (what hasn't yet been covered by COMP11120), 1 and 2.1

²Note that this means that work cannot be submitted late for this unit

J.E. Hopcroft, R. Motwani, and J.D. Ullman. **Introduction to Automata Theory, Languages, and Computation.** *Addison Wesley*, second edition, 2001. ISBN 0-201-44124-1.

This account is derived from the classical textbook on the subject. It is more verbose than Sipser, and it aims to include up-to-date examples and applications, and it develops the material in much detail. A number of illustrated examples are also included. I have heard that this book is very popular among students. 2006/7 editions are also available. About £50. Relevant for this course are Chapters 1 to 5.

Get in touch

Feedback. I'm always keen to get feedback regarding my teaching. Although the course has been taught a few times the notes probably still contain some errors. I'd like to hear about any errors that you may find so that I can correct them. Corrections will be available via the course web site. You can talk to me after lectures, send me email (use the address on the title page) or come and talk to me during my Open Office, which is 12:30 on Wednesday. My office is 2.52 in the Kilburn Building.

Reward! If you can find a substantial mistake in the lecture notes (that is, more than just a typo, or a minor language mistake) you get a chocolate bar. Badges may be awarded for spotting mistakes or making other interesting contributions to the course unit.

Wanted. I'm on the lookout for good examples to use, either in the notes or in the lectures. These should be interesting, touch important applications of the material, or be fun. Again, rewards may be available for a really good reasonably substantial example.

Acknowledgements. First of all, many thanks to Andrea Schalk who taught this unit in previous years and provided this comprehensive set of notes. Many others have contributed to the development of the material, including Giles Reger, Howard Barringer, Pete Jinks, Djihed Affi, Francisco Lobo, Andy Ellyard, Ian Worthington, Peter Sutton, James Bedford, Matt Kelly, Cong Jiang, Mohammed Sabbar, Jonas Lorenz, Tomas Markevicius, Joe Razavi, Will Brown and Rashmica Gupta. Thanks also to students of the unit who have identified errors and provided suggestions for improvement.

Contents

Organisation	1
1 Introduction	10
2 Describing languages to a computer	12
2.1 Terminology	12
2.2 Defining new languages from old ones	13
2.3 Describing languages through patterns	14
2.4 Regular expressions	16
2.5 Matching a regular expression	17
2.6 The language described by a regular expression	18
2.7 Regular languages	20
2.8 Summary	21
3 How do we come up with patterns?	22
3.1 Using pictures	22
3.2 Following a word	24
3.3 Finite state automata	26
3.4 Non-deterministic automata	28
3.5 Deterministic <i>versus</i> non-deterministic	32
3.5.1 Using Automata	38
3.5.2 Describing Languages	38
3.6 From automata to patterns	39
3.6.1 An Alternative Approach: GNFA's	48
3.7 From patterns to automata	49
3.8 Properties of regular languages	58
3.9 Equivalence of Automata	65
3.10 Limitations of regular languages	74
3.11 Summary	76
4 Describing more complicated languages	77
4.1 Generating words	77
4.2 Context-free grammars	78
4.3 Regular languages are context-free	82
4.4 Parsing and ambiguity	83
4.5 A programming language	86
4.6 The Backus-Naur form	87
4.7 Properties of context-free languages	88
4.8 Limitations of context-free languages	89
4.9 Summary	89

A	A Mathematical Approach	91
A.1	The basic concepts	91
A.2	Regular expressions	94
A.3	Finite state automata	95
A.4	The Pumping Lemma	99
	A.4.1 Using the Pumping Lemma	100
A.5	Grammars	101
	Glossary	103

Chapter 1

Introduction

Computers need to be able to interpret input from a keyboard. They have to find commands and then carry out appropriate actions. With a command line interface this may appear a fairly simple problem, but, of course, the keyboard is not the only way of feeding a computer with instructions. Computers also have to be able to read and interpret files, for example when it comes to compiling a program from a language such as **Java**, **C** or **ML** and they then have to be able to run the resulting code. In order to do this computers have to be able to split the input into strings, parse these strings, and turn those into instructions (such as Java bytecode) that the machine can carry out. How does a computer organize potentially quite large input? How does it decide what variables to create, and what values to give these? In this part of *Fundamentals of Computation* we look at some techniques used for this purpose.

In order to organize (or parse) the given text-like input a computer has to be able to recognize specific strings, for example simple commands (such as **if** or **else**). The ability to find certain strings is useful in other contexts as well. When a search engine such as Google looks for a string or phrase entered by a user it certainly has to be capable of telling when it has found the string in question. However, most search engines do their job in a much more clever way than that: They will also recognize plurals of strings entered, and if such a string should be a verb of the English language it will typically recognize different forms, such as *type*, *typing* and *typed*. How can we come up with clever ways of searching for something more than just a given fixed string? Again this is an issue that this course covers.

How does the web server for the Intranet in the School of Computer Science check whether you are accessing it from within the University?

How do online shopping assistants work? They find a given product at various shops and tell you what it costs. How do they extract the information, in particular since some of the pages describing the product are created on demand only?

Imagine you had to write a piece of code that accepts a string as an input and checks whether this string could be a valid email address.

Imagine you had to write a program that goes through a number of emails in a folder and returns all the subject lines for those (every email client has a way of doing this, of course). How would you do that?

Imagine you had to write a program that goes through a bunch of documentation files and checks whether they contain any double typed words (such as ‘the the’). This is a very common mistake, and a professional organization would want to remove these before delivering the software to a customer.

What if

- Your program had to work across lines (and an arbitrary number of spaces)?
- You had to take into account that the first doubled word might be capitalized?
- The two doubled words might be separated by html tags?

Imagine you had to write code for a vending machine that takes coins and in return delivers variously priced goods. How does the machine keep track of how much money the user has put in so far, and how to produce correct change from the coins it holds? It's your job to keep track of all the cases that may arise.

Imagine you had to write a piece of code that takes as input a file containing a **Java** program and checks whether the curly brackets `{ }` are all properly balanced.

This course is going to help you with understanding how such tasks can be carried out. However, we won't spend much time writing code. The tools we describe here could be used when using the Unix **grep** command, or when programming in languages such as **Perl**, **Python**, **Tcl**, or even when using the GNU Emacs editor. It is more important to me that you understand the ideas, and how to express them in a reasonably formal way (so that programming them then becomes easy).

The material can be split into three parts, dealing with the following issues.

- How do we describe to a computer what strings to look for?
- How can we come up with the right description to solve our problem?
- How can we write code that checks whether some other piece of code is correctly formed, whether the opening and closing brackets match, or whether a webpage is written in valid html?

Chapter 2

Describing languages to a computer

In order to solve the kinds of problems that are mentioned in Chapter 1 we need to be able to describe to a computer what it is looking for.

Only in very simple cases will this consist of just one or two strings—in general, we want the computer to look for a much larger collection of words. This could be the set of all possible IP addresses within the University, or it could be the collection of all strings of the form ‘Subject: . . .’ (to pick out emails on a particular topic), or it could be the set of all strings of the form: $s \sqcup s$ in the simplest case of finding all doubled words—this one won’t take care of doubled words spread over two lines, nor of the first word being capitalized.

2.1 Terminology

In order to talk about how we do this in practice we have to introduce a few terms so that we can talk about the various concepts involved. You can find formal definitions of these Appendix A.

- A **symbol** is a building block for a string. Symbols cannot be sub-divided, they are the atoms of everything we build. In the theory of formal languages they are usually called **letters**. Examples: $a, A, 0, 1, \%, @$. We use letters from the end of the Roman alphabet, such as x, y, z to refer to an arbitrary symbol.
- An **alphabet** is a collection of letters. We think of it as a set. At any given time we are only interested in those words we can build from a previously specified alphabet. Examples: $\{0, 1\}, \{0, 1, \dots, 9\}, \{a, b, \dots, z\}, \{a, b, \dots, z, A, B, \dots, Z\}$. We use capital Greek letters, typically Σ , to refer to an arbitrary alphabet.
- A **string** is something we build from 0 or more symbols. In the theory of formal languages we usually speak of **words**. Examples: $ababab, 0, 1001$. Note that every letter can be viewed as a one-letter word. We use letters s, t, u to refer to an arbitrary word.
- The **empty word** (which consists of 0 letters) is a bit difficult to denote. We use ϵ (the Greek letter ‘ ϵ ’) for this purpose.
- **Concatenation** is the operation we perform on words (or letters) to obtain longer words. When concatenating a with b we get ab , and we can concatenate that with the word ba to obtain $abba$. If we concatenate 0 letters we get the word ϵ . When we concatenate any word with the word ϵ we obtain the same word.
- We use the notation of **powers** for concatenation as follows: If s is a word then $(s)^n$ is the word we get by concatenating n copies of s . For example, $(010)^3 = 010010010$, $1^2 = 11$, $b^0 = \epsilon$ and $c^1 = c$.
- A **language** is a collection of words, which we think of as a set. Examples are $\{\epsilon\}$, \emptyset , $\{ab, abc, aba\}$ and $\{a^n \mid n \in \mathbb{N}\}$. We use letters such as $\mathcal{L}, \mathcal{L}_1$ and \mathcal{L}' to refer to an arbitrary language.

In these notes we are interested in how we can describe languages in various ways.

If a language is finite then we can describe it quite easily: We just have to list all its elements. However, this method fails when the language is infinite, and even when the language in question is very large. If we want to communicate to a computer that it is to find all words of such a language we have to find a concise description.

2.2 Defining new languages from old ones

One way of describing languages is using set-theoretic notation (see Appendix A for more detail). In the main development here we try to avoid being overly mathematical, but there are some operations on languages we need to consider. There are seven of them:

- **Union.** Since languages are just sets we can form their unions.
- **Intersection.** Since languages are merely sets we can form their intersections.
- **Set difference.** If \mathcal{L}_1 and \mathcal{L}_2 are languages we can form

$$\mathcal{L}_1 \setminus \mathcal{L}_2 = \{s \in \mathcal{L}_1 \mid s \notin \mathcal{L}_2\}.$$

- **Complement.** If \mathcal{L} is the language of all words over some alphabet Σ , and \mathcal{L}' is a subset of \mathcal{L} then the complement of \mathcal{L}' in \mathcal{L} is $\mathcal{L} \setminus \mathcal{L}'$, the set of all words over Σ which are not contained in \mathcal{L}' .
- **Concatenation.** Because we can concatenate words we can use the concatenation operation to define new languages from existing ones by extending this operation to apply to languages as follows. Let \mathcal{L}_1 and \mathcal{L}_2 be languages over some alphabet Σ . Then

$$\mathcal{L}_1 \cdot \mathcal{L}_2 = \{s \cdot t \mid s \in \mathcal{L}_1 \text{ and } t \in \mathcal{L}_2\}.$$

- **n -ary Concatenation.** If we apply concatenation to the same language by forming $\mathcal{L} \cdot \mathcal{L}$ there is no reason to stop after just one concatenation. For an arbitrary language \mathcal{L} we define

$$\mathcal{L}^n = \{s_1 \cdot s_2 \cdot \dots \cdot s_n \mid s_i \in \mathcal{L} \text{ for all } 1 \leq i \leq n\}.$$

We look at the special case

$$\mathcal{L}^0 = \{s_1 \cdot \dots \cdot s_n \mid s_i \in \mathcal{L} \text{ for all } 1 \leq i \leq n\} = \{\epsilon\},$$

since ϵ is what we get when concatenating 0 times.

- **Kleene star.** Sometimes we want to allow any (finite) number of concatenations of words from some language. The operation that does this is known as the **Kleene star** and written as $(-)^*$. The definition is:

$$\begin{aligned} \mathcal{L}^* &= \{s_1 s_2 \dots s_n \mid n \in \mathbb{N}, s_1, s_2, \dots, s_n \in \mathcal{L}\} \\ &= \bigcup_{n \in \mathbb{N}} \mathcal{L}^n. \end{aligned}$$

Note that

$$\emptyset^* = \bigcup_{n \in \mathbb{N}} \emptyset^n = \{\epsilon\} = \mathcal{L}^0$$

for all \mathcal{L} , which may be a bit unexpected.

Exercise 1. *Language Concatenation*

To gain a better understanding of the operation of concatenation on languages, carry out the following tasks.

- (a) Calculate $\{a, aaa, aaaaaa\} \cdot \{\epsilon, bb, bbb\}$.
- (b) Calculate $\{\epsilon, a, aa\} \cdot \{aa, aaa\}$.
- (c) Calculate $\{a, a^3, a^6\} \cdot \{b^0, b^2, b^3\}$.
- (d) Describe the language $\{0^m 1^n \mid m, n \in \mathbb{N}\}$ as the concatenation of two languages.
- (e) Calculate $\{0, 01, 001\}^2$.

Exercise 2. *Kleene Star*

To gain a better understanding of the Kleene star operation carry out the following tasks. You may use any way you like to describe the infinite languages involved: Just using plain English is the easiest option. I would prefer not to see ... in the description.

- (a) Calculate $\{a, b\}^*$, $\{a\}^* \cup \{b\}^*$, $\{a\}^* \cap \{b\}^*$, $\{aa\}^* \setminus \{aaaa\}^*$ and the complement of $\{a\}^*$ in the language of all words over the alphabet $\{a, b\}$.
- (b) Calculate $\{0^{2n} \mid n \in \mathbb{N}\}^*$.
- (c) Describe the set of all words built from a given alphabet Σ using the Kleene star operation.

2.3 Describing languages through patterns

By definition a language is a set. While sets are something mathematicians can deal with very well, and something that computer scientists have to be able to cope with, computers aren't at all well suited to make sense of expressions such as

$$\{(01)^n \mid n \in \mathbb{N}\}.$$

There are better ways of showing a computer what all the words are that we mean in a particular case. We have already seen that we can express the language \mathcal{L} in a different way, namely as

$$\{01\}^*.$$

If there is only one word in a language, here $\{01\}$, we might as well leave out the curly brackets and write¹ $(01)^*$. That is actually something a machine can cope with, although we would have to use a slightly different alphabet, say $(01)^*\wedge$. All a computer now has to do is to compare: Does the first character² of my string equal 0? Does the next one³ equal 1? And so on.

What we have done here is to create a *pattern*. It consists of various characters of the alphabet, concatenated. We are allowed to apply the Kleene star to any part of the string so created, and we use brackets $()$ to indicate which part should be affected by the star. A computer can then *match* this pattern.

Are these all the patterns we need? Not quite. How, for example, would we describe the language

$$\{0^n \mid n \in \mathbb{N}\} \cup \{1^n \mid n \in \mathbb{N}\} = \{x^n \mid x = 0 \text{ or } x = 1\}?$$

¹Can you see why we need brackets here?

²What should the computer do if the string is empty?

³What if there isn't a next symbol?

We can't use either of 0^*1^* or (arguably worse) $(01)^*$ because both of these include words that contain 0 as well as 1, whereas any word in our target language consists entirely of 0s or entirely of 1s. We need to have a way of saying that either of two possibilities might hold. For this, we use the symbol $|$. Then we can use $0^*|1^*$ to describe the language above.

Exercise 3. *Pattern Matching*

Which of the following words match the given patterns?

Pattern	ϵ	a	ab	b	aba	$abab$	aab	$aabb$	aa
$(ab)^*$									
a^*b^*									
$(a b)$									
$(a b)^*$									
$ab b a^*$									

Exercise 4. *Pattern Matching*

Describe all the words matching the following patterns. For finite languages just list the elements, for infinite ones, try to describe the words in question using English. If you want to practise using set-theoretic notation, add a description in that format.

- (a) $0|1|2$
- (b) 0^*1 ,
- (c) 0^*1^* ,
- (d) $(01)^*0^*$,
- (e) $(01)^*(01)^*$.
- (f) $0|1^*|2$

Exercise 5. *Pattern Matching*

Describe all the words matching the following patterns. For finite languages just list the elements, for infinite ones, try to describe the words in question using English. If you want to practise using set-theoretic notation, add a description in that format.

- (a) $(0|1)(0|2)2$
- (b) $(01|10)$
- (c) $(01)^*1$,
- (d) $(010)^*$,
- (e) $(0|1)^*$
- (f) $0^*|1^*|2^*$

Exercise 6. Pattern Matching

Describe all the words matching the following patterns. For finite languages just list the elements, for infinite ones, try to describe the words in question using English. If you want to practise using set-theoretic notation, add a description in that format.

- (a) $a|b|c$
- (b) 01^* ,
- (c) 1^*0^* ,
- (d) $(10)^*0^*$,
- (e) $(10)^*(10)^*$.
- (f) $a|b|c^*$

2.4 Regular expressions

So far we have been using the idea of a pattern intuitively—we have not said how exactly we can form patterns, nor have we properly defined when a word matches a pattern. It is time to become rigorous about these issues.

For reason of completeness we will need two patterns which seem a bit weird, namely ϵ (the pattern which is matched precisely by the empty word ϵ) and \emptyset (the pattern which is matched by no word at all).

Definition 1. Let Σ be an alphabet. A **pattern** or **regular expression** over Σ is any word over

$$\Sigma^{\text{pat}} = \Sigma \cup \{\emptyset, \epsilon, |, *, (,)\}$$

generated by the following recursive definition.

Empty pattern The character \emptyset is a pattern;

Empty word the character ϵ is a pattern;

Letters every letter from Σ is a pattern;

Concatenation if p_1 and p_2 are patterns then so is (p_1p_2) ;

Alternative if p_1 and p_2 are patterns then so is $(p_1|p_2)$;

Kleene star if p is a pattern then so is (p^*) .

In other words we have defined a *language*⁴, namely the language of all regular expressions, or patterns. Note that while we are interested in words over the alphabet Σ we need *additional symbols* to create our patterns. That is why we have to extend the alphabet to Σ^{pat} .

In practice we will often leave out some of the brackets that appear in the formal definition—but only those brackets that can be uniquely reconstructed. Otherwise we would have to write $((0|1)^*0)$ instead of the simpler $(0|1)^*0$. In order to be able to do that we have to define how to put the brackets back into such an expression. We first put brackets around any occurrence of $*$ with the sub-pattern immediately to its left, then around any occurrence of concatenation, and lastly around the alternative operator $|$. Note that every regular expression with all its brackets

⁴In Chapter 4 we look at how to describe a language like that—it cannot be done using a pattern.

has precisely one way of building it from the rules—we say that patterns are *uniquely parsed*⁵. This isn't quite true once we have removed the brackets: $(0|(1|2))$ turns into $0|1|2$ as does $((0|1)|2)$. However, given the way we use regular expressions this does not cause any problems. Note that many computer languages that use regular expressions have additional operators for these (see Exercise 9). However, these exist only for the convenience of the programmer and don't actually make these regular expressions more powerful. We say that they have the same *power of expressivity*. Whenever I ask you to create a regular expression or pattern it is Definition 1 I expect you to follow.

2.5 Matching a regular expression

This language is **recursively defined**: There are base cases consisting of the \emptyset , ϵ and all the letters of the underlying alphabet, and three ways of constructing new patterns from old ones. For those of you who are taking COMP11120, you will also see this kind of definition used for defining, for example lists.⁶ As we will see later on, the use of a recursive definition allows us to apply an inductive argument when proving theorems about our definitions.

If we have to define something for all patterns we may now do so by doing the following:

- Say how the definition works for each of the base cases;
- assuming we have defined the concept for p , p_1 and p_2 , say how to obtain a definition for
 - p_1p_2 ,
 - $p_1|p_2$ and
 - p^* .

Note that in the definition of a pattern no meaning is attached to any of the operators, or even the symbols that appear in the base cases. We only find out what the intended meaning of these is when we define when a word matches a pattern. This is a common approach in Computer Science: we define a *syntax* for a language or system which tells us how to write down or describe expressions, and then provide a *semantics* which tells us how to interpret those expressions. We define matching by giving a recursive definition as outlined above.

Definition 2. Let p be a pattern over an alphabet Σ and let s be a word over Σ . We say that s *matches* p if one of the following cases holds:

Empty word the pattern is ϵ and s is the empty word ϵ ;

Base case the pattern $p = x$ for a character x from Σ and $s = x$;

Concatenation the pattern p is a concatenation $p = (p_1p_2)$ and there are words s_1 and s_2 such that s_1 matches p_1 , s_2 matches p_2 and s is the concatenation of s_1 and s_2 ;

Alternative the pattern p is an alternative $p = (p_1|p_2)$ and s matches p_1 or p_2 (it is allowed to match both);

Kleene star the pattern p is of the form $p = (q^*)$ and s can be written as a finite concatenation $s = s_1s_2 \cdots s_n$ such that s_1, s_2, \dots, s_n all match q ; this

⁵We will return to the issue of parsing expressions and ambiguity in Chapter 4.

⁶Students who are on the joint honours CS and Maths programme don't take COMP11120, but they should be able to grasp these ideas without problems.

includes the case where s is empty (and thus an empty concatenation, with $n = 0$).

Note that there is no word matching the pattern \emptyset .

Exercise 7. *Pattern Matching*

Calculate all the words matching the patterns ϵ and a (for $a \in \Sigma$) respectively.

Exercise 8. *Pattern Matching Words*

For the given pattern, and the given word, employ the recursive Definition 2 to demonstrate that the word does indeed match the pattern:

- (a) the pattern $(ab)^*a$ and the word aba .
- (b) the pattern $(0|1)^*10$ and the word 10010 ,
- (c) the pattern $(0^*|1^*)10$ and the word 0010 ,
- (d) the pattern $(abc)^*a$ and the word a .

Exercise 9. *Linux Tools*

Here are some examples of the usage of regular expressions in ‘the real world’.

- (a) Print out the manual page for the Linux command `grep` (type `man grep` to get that page). Now argue that for every regular expression understood by `grep` there is a regular expression according to Definition 1 that has precisely the same words matching it.^a
- (b) Give a command-line instruction which will show all the subjects and authors of mails contained in some directory. Make sure to test your suggested answer on a computer running Linux. *Hint: Try using `egrep`.*
- (c) Give a regular expression that matches all IP addresses owned by the University of Manchester. *You probably want to use some kind of shortcut notation.*

^aTheoreticians want their patterns with as few cases as possible so as to have fewer cases for proofs by induction. Practitioners want lots of pre-defined shortcuts for ease of use. This exercise shows that it doesn’t really matter which version you use.

2.6 The language described by a regular expression

Given a pattern we can now define a language based on it.

Definition 3. Let p be a regular expression over an alphabet Σ . The **language defined by pattern** p , $\mathcal{L}(p)$ is the set of all words over Σ that match p . In other words

$$\mathcal{L}(p) = \{s \in \Sigma^* \mid s \text{ matches } p\}.$$

Note that different patterns may define the same language, for example $\mathcal{L}(0^*) = \mathcal{L}(\epsilon|00^*)$.

Note that we can also define the language given by a pattern in a different way, namely recursively.

- $\mathcal{L}(\emptyset) = \emptyset$;

- $\mathcal{L}(\epsilon) = \{\epsilon\}$;
- $\mathcal{L}(x) = \{x\}$ for $x \in \Sigma$;
- and for the operations
 - $\mathcal{L}(p_1 p_2) = \mathcal{L}(p_1) \cdot \mathcal{L}(p_2)$;
 - $\mathcal{L}(p_1 | p_2) = \mathcal{L}(p_1) \cup \mathcal{L}(p_2)$;
 - $\mathcal{L}(p^*) = (\mathcal{L}(p))^*$.

We can use this description to calculate the language defined by a pattern as in the following example.

$$\begin{aligned}
 \mathcal{L}((0|1)^*) &= (\mathcal{L}(0|1))^* \\
 &= (\mathcal{L}(0) \cup \mathcal{L}(1))^* \\
 &= (\{0\} \cup \{1\})^* \\
 &= \{0, 1\}^*
 \end{aligned}$$

This is the language of all words over the alphabet $\{0, 1\}$.

Exercise 10. *Defining Languages*

Use this recursive definition to find the languages defined by the following patterns, that is the languages $\mathcal{L}(p)$ for the following p .

- (a) $(0|1)^*$
- (b) $(0^*|1^*)$
- (c) $(01)^*0$
- (d) $(00)^*$
- (e) $((0|1)(0|1))^*$.

Can you describe these languages in English?

In order to tell a computer that we want it to look for words belonging to a particular language we have to find a pattern that describes it. The following exercise lets you practise this skill.

Exercise 11. *Finding Patterns*

Find a regular expression p over the alphabet $\{0, 1\}$ such that the language defined by p is the one given. *Hint: For some of the exercises it may help not to think of a pattern that is somehow formed like the strings you want to capture, but to realize that as long as there's one way of matching the pattern for such a string, that's good enough.*

- (a) All words that begin with 0 or end with 1.
- (b) All words whose length is at most 4.
- (c) All words that aren't equal to 11 or 111.

Exercise 12. *Finding Patterns*

Repeat Exercise 11 for the following languages.

- (a) All words that contain at least three 0s.
- (b) All words which contain the string 101, that is, a 1 followed by a 0 and then a 1.
- (c) All words whose length is at least 3.

Exercise 13. *Finding Patterns*

Repeat Exercise 11 for the following languages^a.

- (a) All words for which every letter at an even position is 0.
- (b) All words that aren't equal to the empty word.
- (c) All words that contain an even number of 0.
- (d) All words whose number of 0s is divisible by 3.
- (e) All words that do not contain the string 11.
- (f) All words that do not contain the string 101.
- (g) All words that contain an even number of 0s and whose number of 1s is divisible by 3.

^aIf you find the last few of these really hard then skip them for now. The tools of the next chapter should help you finish them.

Exercise 14. *Finding Patterns*

Find a regular expression p over the alphabet $\{a, b, c\}$ such that the language defined by p is the one given.

- (a) All the words that do not contain the string ab .
- (b) All the words that don't contain the letter b .

Exercise 15. *Finding Patterns*

Repeat Exercise 14 for the following.

- (a) All the words that do not contain the string aba .
- (b) All the words where every a is immediately followed by b .

2.7 Regular languages

We have now described a certain category of languages, namely those that can be defined using a regular expression. It turns out that these languages are quite important so we give them a name.

Definition 4. A language L is **regular** if it is the set of all words matching some regular expression, that is, if there is a pattern p such that $\mathcal{L} = \mathcal{L}(p)$.

Regular languages are not rare, and there are ways of building them. Nonetheless in Chapter 4 we see that not all languages of interest are regular.

Proposition 2.1. Assume that \mathcal{L} , \mathcal{L}_1 , and \mathcal{L}_2 are regular languages over an alphabet Σ . Then the following also are regular languages.

- (a) $\mathcal{L}_1 \cup \mathcal{L}_2$
- (b) $\mathcal{L}_1 \cdot \mathcal{L}_2$
- (c) \mathcal{L}^n
- (d) \mathcal{L}^*

Proof of this follows easily by considering the regular expressions that would define the languages give. For example, if \mathcal{L}_1 and \mathcal{L}_2 are regular, then there must be regular expressions e_1 and e_2 that match precisely the words in \mathcal{L}_1 and \mathcal{L}_2 respectively. The regular expression $e_1|e_2$ will then match precisely those words in $\mathcal{L}_1 \cup \mathcal{L}_2$ and the language is thus regular. Proof of the remaining claims in the Proposition is left as an exercise for the reader. Section 3.8 considers some additional properties of regular languages.

Regular expressions are very handy when it comes to communicating to a computer quite complicated collections of words that we might be looking for. However, coming up with a regular expression that describes precisely those words we have in mind isn't always easy.

The definitions that we have introduced in this chapter tell us what it means for a word to match a regular expression. What we don't explicitly have here though is an easy way of actually *doing* this. Could you write a programme that takes a regular expression and a word and tells us whether or not there is a match? What would be the problems that you might face when trying to do this? In the next chapter we look at different ways of describing languages using *automata*. Automata can make it easier for human beings to ensure that they are describing the language they had in mind and they are relatively straightforward to implement. We also introduce algorithms that translate between automata and patterns.

2.8 Summary

- A *language* is a set of *words* over an *alphabet* of *symbols* or *letters*.
- In order to describe a language to a computer we can use *regular expressions*, also known as *patterns*.
- For regular expressions we have the notion of it being *matched* by some particular word.
- Each regular expression defines a language, namely the set of all words that match it.
- We say that a language is *regular* if there is a pattern that defines it.

Chapter 3

How do we come up with patterns?


As you should have noticed by now (if you have done the last two exercises from the previous section) coming up with a pattern that describes a particular language can be quite difficult. One has to develop an intuition about how to think of the words in the desired language, and turn that into a characteristic for which a pattern can be created. While regular expressions give us a format that computers understand well, human beings do much better with other ways of addressing these issues.

3.1 Using pictures

Imagine somebody asked you to check a word of an unknown length to see whether it contains an even number of 0s. The word is produced one letter at a time. How would you go about that?


In all likelihood you would not bother to count the number of 0s, but instead you would want to remember whether you have so far seen an odd, or an even, number of 0s. Like flicking a switch, when you see the first 0, you'd remember 'odd', when you see the second, you'd switch back (since that's the state you'd start with) to 'even', and so forth.

If one wanted to produce a description of this idea it might look like this:




Every time we see 0 we switch from the 'even' state to the 'odd' state and *vice versa*.


If we wanted to give this as a description to somebody else then maybe we should also say what we are doing when we see a letter other than 0, namely stay in whatever state we're in. Let's assume we are talking about words consisting of 0s and 1s. Also, we'd like to use circles for our states because they look nicer, so we'll abbreviate their names.



So now somebody else using our picture would know what to do if the next letter is 0, and what to do if it is 1. But how would somebody else know where to begin?




We give them an arrow that points at the state one should start in. However, they would still only know whether they finished in the state called *E* or the one called *O*, which wouldn't tell them whether this was the desired outcome or not. Hence we mark the state we want to be in when the word comes to an end. Now we do have a complete description of our task.




We use a double circle for a state to indicate that if we end up there then the word we looked at satisfied our criteria.

Let's try this idea on a different problem. Let's assume we're interested in whether our word has 0 in every position that is a multiple of three, that is, the third, sixth, ninth, *etc*, letters, if they exist, are 0.


So we start in some state, and we don't care whether the first letter is 0 or 1, but we must remember that we've seen one letter so that we can tell when we have reached the first letter, so we'd draw something like this:




Similarly for the second letter.




Now the third. This time something happens: If we see 0, we're still okay, but if we see 1 then we need to reject the word.



So what now? Well, if we reach the state labelled 4 then no matter what we see next, we will never accept the word in question of satisfying our condition. So we simply think of this state as one that means we won't accept a word that ends there. All the other states are fine—if we stop in any of them when reading a word it does satisfy our condition. So we mark all the other states as good ones to end in, and add the transitions that keep us in state 4.



But what if the word is okay until state 3? Then we have to start all over again, not caring about the next letter or the one after, but requiring the third one to be 0. In other words, we're in the same position as at the start—so the easiest thing to do is *not to create state 3*, but instead to have that edge go back to state 0.




Exercise 16. *Drawing Automata*

For the languages described in parts (a) and (b) of Exercise 12, draw a picture as in the examples just given.


3.2 Following a word

Consider the following picture.




What happens when we try to follow a word, say *abaa*? We also sometimes refer to an automaton as *consuming* a word.


To begin with we are in the start state with the word *abaa*.




Our word is *abaa*. We see *a*, so we follow the edge labelled *a*.




Having followed the first letter, we can drop it, and our word is now baa .
 Now we see the letter b , so we go once round the loop in our current state.



Again, we drop the letter we've just dealt with, leaving us with aa .
 Now we have aa so we follow the edge labelled a



Lastly we have the word a , so we follow the edge to the right, labelled a, b .



We end up in a non-accepting state so this *wasn't* one of the words we were looking for.

Exercise 17. *Following Words*

Follow these words through the above automaton. If you can, try to describe the words which end up in the one desirable state (the double-ringed one).


- (a) *abba*
- (b) *baba*
- (c) *baab*.

3.3 Finite state automata

So far we have been very informal with our pictures. It is impossible to define an automaton without becoming fairly mathematical. For any of these pictures we must know which *alphabet* Σ we are considering.

Every automaton consists of

- a number of states,
- one state that is *initial*,
- a number (possibly 0) of states that are *accepting*
- for every state and every letter x from the alphabet Σ precisely one *transition* from that state to another labelled with x .



Formally we may consider the set of all the states in the automaton, say Q . One of these states is the one we want to start in, called the *start state* or the *initial state*. Some of the states are the ones that tell us if we end up there we have found the kind of word we were looking for. We call these *accepting states*. They form a subset, say F , of Q .

The edges in the graph are a nice way of visualizing the transitions. Formally what we need is a function that

- takes as its input
 - a state and
 - a letter from Σ
- and returns

- a state.

We call this the *transition function*, δ . It takes as inputs a state and a letter, so the input is a pair (q, x) , where $q \in Q$ and $x \in \Sigma$. That means that the input comes from the set

$$Q \times \Sigma = \{(q, x) \mid q \in Q, x \in \Sigma\}.$$

Its output is a state, that is an element of Q . So we have that

$$\delta : Q \times \Sigma \longrightarrow Q.$$

The formal definition then is as follows.

Definition 5. Let Σ be a finite alphabet of symbols. A **(deterministic) finite^a automaton** or **DFA**, over Σ consists of the following:

- A finite non-empty set Q of states;
- a particular element of Q called the start state (which we often denote with q_\bullet);
- a subset F of Q consisting of the accepting states;
- a transition function δ which for every state $q \in Q$ and every symbol $x \in \Sigma$ returns the next state $\delta(q, x) \in Q$, so δ is a function from $Q \times \Sigma$ to Q . When $\delta(q, x) = q'$ we often write

$$q \xrightarrow{x} q'.$$

We sometimes put these four items together in a quadruple $(Q, q_\bullet, F, \delta)$.

^aAll our automata have a finite number of states and we often drop the word ‘finite’ when referring to them in these notes.

Sometimes people also refer to a **finite state machine**.

Note that for every particular word there is precisely *one path* through the automaton: We start in the start state, and then read off the letters one by one. The transition function makes sure that we will have precisely one edge to follow for each letter. When we have followed the last letter of the word we can read off whether we want to accept it (if we are in an accepting state) or not (otherwise). That’s why these automata are called *deterministic*; we see non-deterministic automata below.

The deterministic nature of the automata means that they are (relatively) easy to implement. Whenever we see a letter the transition function tells us which state to move to. There is no choice, and a single state is active at any one time. Exercise 107 looks a little closer at this question.

For every word $x_1x_2\cdots x_n$ we have a uniquely determined sequence of states q_0, q_1, \dots, q_n such that $q_0 = q_\bullet$ and

$$q_0 \xrightarrow{x_1} q_1 \xrightarrow{x_2} q_2 \longrightarrow \cdots \xrightarrow{x_n} q_n.$$

We *accept* the word if and only if the last state reached, q_n is an accepting state. Formally it is easiest to define this condition as follows.

Definition 6. A word $s = x_1\cdots x_n$ over Σ is **accepted by the deterministic finite automaton** $(Q, q_\bullet, F, \delta)$ if

- $\delta(q_\bullet, x_1) = q_1$,
- $\delta(q_1, x_2) = q_2$,

- $\dots,$
- $\delta(q_{n-1}, x_n) = q_n$ and
- $q_n \in F$, that is, q_n is an accepting state.

In particular, the empty word is accepted if and only if the start state is an accepting state. Just as was the case for regular expressions we can view a DFA as defining a language.

Definition 7. We say that a *language is recognized by a finite automaton* if it is the set of all words accepted by the automaton.

Exercise 18. *DFA's for Patterns*

Design DFA's that accept precisely those words given by the languages described in the various parts of Exercise 11.

Exercise 19. *DFA's for Patterns*

Design a DFA that accepts precisely those words given by the language described in part (c) of Exercise 12.

Exercise 20. *DFA's for Patterns*

Design DFA's that accept precisely those words given by the languages described in Exercise 13.

Exercise 21. *DFA's for Patterns*

Design DFA's that accept precisely those words given by the languages described in Exercise 14.

Exercise 22. *DFA's for Patterns*

Design DFA's that accept precisely those words given by the languages described in Exercise 15.


3.4 Non-deterministic automata

Sometimes it can be easier to draw an automaton that is *non-deterministic*. What this means is that from some state, there may be several edges labelled with the *same letter*. As a result, there is then no longer a unique path when we follow a particular word through the automaton. Hence the procedure of following a word through an automaton becomes more complicated—we have to consider a number of possible paths we might take.


Consider an automaton that accepts words over the alphabet $\{0, 1\}$ that start with a 0 and have a 1 as the second last character.¹ A regular expression for this language is relatively easy to produce: $0(0|1)^*1(0|1)$. But the automaton is less straight forward.

We know that any word we accept has to start with 0, so if the first letter is a 1 we will not accept the word, so we start with something like this:

¹Can you think of any properties that can we immediately determine that words in this language will have? For example, all words must have at least three letters. Why?




Now if we see any number of 0s we know we're not going to accept the word. But if we see a 1, then we will accept the word if there is just one more symbol after that (and no more). So we might draw something like the following:



But this isn't quite right. What happens in state E ? If we see another letter then we actually want to be in accepting state *but only if the letter that we saw when moving from state D to state E was a 1*. Things now begin to get rather complex as we find ourselves needing to remember which letter we saw previously.

What you might have found yourself trying to draw is something like this:



This is almost the same as our first attempt, but with one important (and quite subtle) distinction: when we are in state C , and we see a 1, we have a *choice* as to where to go next – either stay at C or move to D .

So if we now follow the word 010 we have two possible paths through the automaton:

From state A , read 0, go to state C .

Read 1 and remain in state C .

Read 0 and remain in state C .

Finish not accepting the word.

From state A , read 0, go to state C .

Read 1 and go to state D .

Read 0 and go to state E .

Finish accepting the word.

We say that the automaton accepts the word if there is *at least one such path* that ends in an accepting state.

So how does the definition of a non-deterministic automaton differ from that of a deterministic one? We still have a set of states Q , a particular start state q_\bullet in Q , and a set of accepting states $F \subseteq Q$.

However, it is no longer the case that for every state and every letter from x there is precisely one edge labelled with x , there may be several. What we no longer have is a transition *function*. Instead we have a transition *relation*.² Given a state q , a letter x and another state q' the relation δ tells us whether or not there is an edge labelled x from q to q' .

Exercise 23. *Deterministic Automata*

Go back and check your solutions to Exercises 18 to 22. Were they all deterministic as required? If not, redo them.

We can turn this idea into a formal definition.

Definition 8. A *non-deterministic finite automaton* or *NFA*, over Σ is given by

- a finite non-empty set Q of states,
- a start state q_\bullet in Q ,
- a subset F of Q of accepting states as well as
- a transition relation δ which relates a pair consisting of a state and a letter in Σ to a state. We often write

$$q \xrightarrow{x} q'$$

if (q, x) is δ -related to q' .

We can now also say when an NFA accepts a word.

Definition 9. A word $s = x_1 \cdots x_n$ over Σ is *accepted by the non-deterministic finite automaton* $(Q, q_\bullet, F, \delta)$ if there are states

$$q_0, q_1, \dots, q_n$$

such that $q_0 = q_\bullet$ and for all $0 \leq i < n$, δ relates (q_i, x_i) to q_{i+1} and such that $q_n \in F$, that is, q_n is an accepting state. The **language recognized by an NFA** is the set of all words it accepts.

An NFA therefore accepts a word $x_1 x_2 \cdots x_n$ if there are states

$$q_\bullet = q_0, q_1, \dots, q_n$$

such that


$$(q_\bullet = q_0) \xrightarrow{x_1} q_1 \xrightarrow{x_2} q_2 \longrightarrow \cdots \xrightarrow{x_n} q_n,$$

and q_n is an accepting state. However, the sequence of states is no longer uniquely determined, and there could potentially be many.

Exercise 24. *NFAs*

Consider the following NFA.


²You will meet relations also in COMP11120. Again, for CM students the concept of relation should be familiar.



Which of the following words are accepted by the automaton? ϵ , *aba*, *bbb*, *bba*, *baa*, *baba*, *babb*, *baaa*, *baaaa*, *baabb*, *baaba*. Can you describe the language consisting of all the words accepted by this automaton?

Exercise 25. NFAs


Consider the following NFA.



Which of the following words are accepted by the automaton? ϵ , *aba*, *bbb*, *bba*, *baa*, *baab*, *baba*, *babb*, *baaba*, *babab*, *baaaa*. Can you describe the language consisting of all the words accepted by this automaton?


Note that in the definition of NFA there is no rule that says that for a given state there *must* be an edge for every label. That means that when following a word through a non-deterministic automaton we might find ourselves stuck in a state because there is no edge out of it for the letter we currently see. If that happens then we know that along our current route, the word will not be accepted by the automaton. While this may be confusing at first sight, it is actually quite convenient. It means that pictures of non-deterministic automata can be quite small.

Take the above example of finding an NFA that accepts all the words that start with 0 and have next to last letter 1. We can give a smaller automaton that does the same job.



You should spend a moment convincing yourself that this automaton does indeed accept precisely the words claimed (that is, the same ones as the previous automaton).³

This is such a useful convention that it is usually also adopted when *drawing deterministic automata*. Consider the problem of designing a DFA that recognizes those words over the alphabet $\{0, 1\}$ of length at least 2 for which the first letter is 0 and the second letter is 1. By concentrating on what it takes to get a word to an accepting state one might well draw something like this:




³Note that unless we have to describe the automaton in another way, or otherwise have reasons to be able to refer to a particular state, there is no reason for giving the states names in the picture.

This is a perfectly good picture of a deterministic finite automaton. However, not all the states, and not all the transitions, are drawn for this automaton: Above we said that for every state, and every letter from the alphabet, there must be a transition from that state labelled with that letter. Here, however, there is no transition labelled 1 from the state 0, and no transition labelled 0 from the state 1.

What *does* the automaton do if it sees 1 in state 0, or 0 in state 1? Well, it discards the word as non-acceptable, in a manner of speaking.


We can complete the above picture to show all required states by assuming there's a hidden state that we may think of as a 'dump'. As soon as we have determined that a particular word can't be accepted we send it off in that dump state (which is certainly not an accepting state), and there's no way out of that state. So all the transitions not shown in the picture above go to that hidden state. With the hidden state drawn our automaton looks like this:



This picture is quite a bit more complicated than the previous one, but both describe the same DFA, and so contain precisely *the same information*. I am perfectly happy for you to draw automata either way when it comes to exam questions or assessed coursework.

Exercise 26. Dump States

Consider the following DFA. Which of its states are dump states? An *unreachable* state (see page 37) is a state that we can never get to when starting at the start state. Which states are unreachable? Draw the simplest automaton recognizing the same language.



Describe the language recognized by the automaton.

Exercise 27. Dump States

Go through the automata you have drawn for Exercise 18 to 22. Identify any dump states in them.

3.5 Deterministic *versus* non-deterministic

So far we have found the following differences between deterministic and non-deterministic automata: For the same problem it is usually easier to design a non-deterministic automaton, and

the resulting automata are often smaller. On the other hand, following a word through a deterministic automaton is straightforward, and so deciding whether the word is accepted is easy. For non-deterministic automata we have to find all the possible paths a word might move along, and decide whether any of them leads to an accepting state. Hence finding the language recognized by an NFA is usually harder than to do the same thing for a DFA of a similar size.


So both have advantages and disadvantages. But how different are they really? Clearly every deterministic automaton can be viewed as a non-deterministic one since it satisfies all the required criteria.

It therefore makes sense to wonder whether there are things we can do with non-deterministic automata that can't be done with deterministic ones. It turns out that this is not the case.


Theorem 3.1. *For every non-deterministic automaton there is a deterministic one that recognizes precisely the same words.*

Algorithm 1, example


Before looking at the general case of Algorithm 1 we consider an example. Consider the following NFA.




We want to construct a deterministic automaton from this step by step. We start with state 0, which we just copy, so it is both initial and an accepting state in our new automaton.




With a , we can go from state 0 to states 1 and 2, so we invent a new state we call 12 (think of it as being a set containing both, state 1 and state 2). Because 2 is an accepting state we make 12 an accepting state too.




With a letter b we can go from state 0 to state 1, so we need a state 1 (think of it as state $\{1\}$) in the new automaton too.




Now we have to consider the states we have just created. In the original automaton from state 1, we can't go anywhere with a , but with b we can go to state 2, so we introduce an accepting state 2 (thought of as $\{2\}$) into our new automaton.




To see where we can go from state 12 we do the following. In the original automaton, with a we can go from 1 nowhere and from 2 to 0 and 2, so we lump these together and say that from 12 we can go to a new state 02 (really $\{0, 2\}$) with a . Because 2 is an accepting state we make 02 one too.




In the original automaton from state 1 with b we can go to state 2, and from state 2 we can go to state 1 with the same letter, so from state 12 we can go back to 12 with b .




In the original automaton with a we can go from state 2 to states 0 and 2, so we need a transition labelled a from state 2 to state 02 in our new DFA.




With b from state 2 we can only go back to 1, so we add this transition to the new automaton.




Now for the new state 02. From 0 we can go with a to states 1 and 2, and from state 2 we can go to states 0 and 2, so taking it all together from state 02 we can go to a new accepting state we call 012.



With b from state 0 we can go to state 1 in the old automaton, and from state 2 we can also only go to state 1 with a b , so we need a transition from the new state 02 to state 1 labelled b .



Following the same idea, from 012 with a we can go back to 012, and with b we can go to 12.




So what we have done here? We have created a new automaton from the old one as follows.

- States of the new automaton are *sets of states* from the old one.

- The start state for the new automaton is the set containing just the initial state from the original automaton.
- A state of the new automaton is an accepting state if one of its states in the old automaton is an accepting state.
- There is a transition labelled x from some state in the new automaton to another if there are states of the old automaton in the two sets that have a corresponding transition.

Exercise 28. NFA to DFA

For the automaton given below carry out the same process. *Hint: If you find this very hard then read on a bit to see if that helps.*



Algorithm 1, general case

We now give the general case.

Assume that we have an NFA $(Q, q_\bullet, F, \delta)$ over some alphabet Σ . Mathematically, we have constructed a new automaton, a DFA this time, as follows:

- The states are given by $\mathcal{P}(Q)$, where \mathcal{P} is the powerset operator, so

$$\mathcal{P}(Q) = \{S \mid S \subseteq Q\}.$$

- The start state is $\{q_\bullet\}$.
- A state S is an accepting state if and only if there is $q \in S$ such that $q \in F$, that is, q is an accepting state in the original automaton. Hence the new set of accepting states is

$$F' = \{S \subseteq Q \mid \exists q \in S \text{ with } q \in F\}.$$

- For a new state $S \in \mathcal{P}(Q)$ and a letter $x \in \Sigma$ the transition labelled with x leads to the state given by the set of all states q' for which there is a $q \in S$ such that there is a transition labelled x from q to q' , that is

$$\{q' \in Q \mid \exists q \in S \text{ with } q \xrightarrow{x} q'\}.$$

In other words, to see where the edge labelled x from S should lead to, go through all the states q in S and collect all those states $q' \in Q$ which have an edge labelled x from q .


Note that if there is no state in S which is labelled with x then $\delta'(S, x)$ is still defined—it is the empty set in that case. Also note that the state given by the empty set is *always* a ‘dump state’. Because there are no states in \emptyset , $\delta(\emptyset, x) = \emptyset$ is true for all $x \in \Sigma$, so it is a state one can never escape from. Also, it cannot be an accepting state because it contains no accepting state as an element.

We call the resulting automaton the **DFA generated by the NFA** $(Q, q_\bullet, F, \delta)$.

For the above example, the new automaton has states \emptyset , $\{0\}$, $\{1\}$, $\{2\}$, $\{0, 1\}$, $\{0, 2\}$, $\{1, 2\}$ and $\{0, 1, 2\}$. In the image above we have kept our labels shorter by only listing the elements of

the corresponding set without any separators or braces. For larger automata this would make it impossible to distinguish between the state $\{12\}$ and the state $\{1, 2\}$, but in such a small example this is harmless.

However, not all the states we have just listed appear in the picture we drew above. This is so because we used a method that only draws those states which are reachable from the start state. We next draw the full automaton as formally defined above.



We already know that we may leave out dump states like \emptyset when drawing an automaton.


We have another extra state in this picture, namely $\{0, 1\}$. This is a state we can never get to when starting at the start state $\{0\}$, so no word will ever reach it either. It is therefore irrelevant when it comes to deciding whether or not a word is accepted by this automaton. We call such states **unreachable** and usually don't bother to draw them.

Exercise 29. Full NFA

For the NFA from Exercise 28, draw the picture of the full automaton with all states (given by the set of all subsets of $\{0, 1, 2\}$, including the unreachable ones) and all transitions.


Exercise 30. NFA to DFA

For following NFA, give a DFA recognizing the same language.




Exercise 31. NFA to DFA

For the following NFA, use *Algorithm 1* to give a DFA recognizing the same language.

**Exercise 32.** *NFA to DFA*

Apply Algorithm 1 to the simplified NFA for language $0(0|1)^*1(0|1)$:



It is possible to avoid drawing unreachable states by doing the following:

- Start with the start state of the new automaton, $\{q_\bullet\}$.
- Pick a state S you have already drawn, and a letter x for which you have not yet drawn a transition out of the state S . Now there must be precisely one state S' to which the transition labelled x leads.
 - If the state S' has already been drawn, draw in the new transition labelled x
 - If the state S' has not been drawn yet, add it to the automaton and draw in the new transition labelled x .

The DFA $(\mathcal{P}(Q), \{q_\bullet\}, F', \delta')$ defined above accepts precisely the same words as the original NFA $(Q, \{q_\bullet\}, F, \delta)$. This establishes Theorem 3.1, which tells us that we can do precisely the same things with NFAs that we can do with DFAs.

3.5.1 Using Automata

Note that we can use these automata also to encode information that at first sight looks quite different from that of recognizing a language!

Exercise 33. *A Vending Machine*

You are asked to design a vending machine that accepts the following coins: 2p, 5p, 10p. Products handled by the machine cost 10p. In order to keep track of when the machine has received 10p (or more), draw a DFA.

3.5.2 Describing Languages

We now have three different ways in which we can describe the same language.

Exercise 34. *Odd Length*

Consider the following language over the alphabet $\{a, b\}$ which consists of all words of odd length. Describe this language in the following ways:

- (a) Using a regular expression.
- (b) Using a DFA.
- (c) As a set. *Unless you work through the Appendix you will probably find this difficult.*

Exercise 35. *Describing Languages*

Consider the following language over the alphabet $\{a, b, c\}$ which consists of all words of odd length that start with a c . Describe this language in the following ways:

- (a) Using a regular expression.
- (b) Using a DFA.
- (c) As a set. *Unless you work through the Appendix you will probably find this difficult.*

3.6 From automata to patterns


We have three ways now of talking about collections of strings, or languages.

- (i) Set-theoretically;
- (ii) Using a regular expression; or
- (iii) Using a (deterministic or non-deterministic) finite automaton

The second of these is particularly good when we have to talk to a computer, and sometimes it's quite easy to come up with the right regular expression. For more complex problems, however, it's often easier (or at least less error-prone) to come up with an automaton. But how do regular expressions and automata connect?

If we have found an automaton for the language we want, how do we communicate it to a computer? In this section we describe an algorithm that takes a DFA and turns it into a pattern describing the same language.

Why do we need an algorithm? Sometimes we can just read off the pattern from the automaton, even if they contain cycles.



Clearly any word that will be accepted has to start with a , can then contain arbitrarily many further as , has a b , and arbitrarily many further bs . That gets the word into an accepting state.

After that, the word may have another a and then it all repeats.

A pattern describing the same language is

$$(aa^*bb^*)(aaa^*bb^*)^*.$$

However, if the automaton is more complicated then reading off a pattern can become very difficult, in particular if there are several accepting states. If the automaton is moreover non-deterministic

the complexity of the task worsens further—usually it is therefore a good idea to first convert an NFA to a DFA using Algorithm 1 from page 36.


In order to show that it is possible to construct a regular expression defining the same language for *every* automaton we have to give an algorithm that works for all automata. This algorithm may look overly complex at first sight, but it really does work for every automaton. If you had to apply it a lot you could do the following:

- Define a data structure of finite automata in whatever language you’re using. In Java you would be creating a suitable class.
- Implement the algorithm so that it takes as input an object of that class, and produces a pattern accordingly.

Algorithm 2, first example

Because the algorithm required is complicated we explain it first using a simple example. With a bit of experience you may be able to read off a pattern from the automaton, but you probably won’t be able to do this with more complicated automata (nor will you be able to program a computer to do the same).

Consider the following DFA.



Before we can apply the algorithm we have to name the states, and the algorithm assumes that the states are, in fact *numbered*. The **order** in which the states are numbered can make a big difference in **how complex** the algorithm is to carry out! Turning an automaton into a regular expression is difficult to do in general since *every loop* in the automaton has to be taken into account. You therefore want to apply the following rules:


- Count the number of states, say there are n . We will number them from $n - 1$ to 0.
- Set $m = n - 1$. Repeat the following until there are no numbered states.
 - Find the state not yet numbered that is involved in the most loops not yet eliminated. Label that state m .
 - $m = m - 1$
 - Remove all those loops from consideration that involve the newly named state.

Or, if you are looking for something simpler:

- Number the states in order of complexity,

that is, the state that has the most paths going through it gets the highest number, the next highest number goes to the next ‘complicated’ state, and so on. If states look roughly similar then it doesn’t matter.

In the example above, the two non-initial states are involved in more paths, but they’re fairly similar, so it doesn’t matter which of them gets the higher number. We use the following numbering.



There are two accepting states. Hence a word that is accepted will either start in state 0 and end in state 0 or start in state 0 and end in state 2. That means the language \mathcal{L} accepted by this automaton is the union of two languages which we write as follows:

$$\mathcal{L} = \mathcal{L}_{0 \rightarrow 0} \cup \mathcal{L}_{0 \rightarrow 2}$$

The indices tell us in which state we start and in which state we finish. This is already a useful observation since we can now concentrate on calculating one language at a time. This is an example of a *divide and conquer* approach, which we often encounter in Computer Science. A problem is broken down into two or more smaller similar subproblems. Each of those subproblems is again broken down until we have problems that are simple enough to be solved directly.

To calculate $\mathcal{L}_{0 \rightarrow 0}$ we note that we can move from state 0 to state 0 in a number of ways. It is so complicated because there are loops in the automaton. What we do is to break up these ways by controlling the use of the state with the highest number, state 2, as follows:

To get from state 0 to state 0 we can

- either not use state 2 at all (that is, go only via states 0 and 1) or
- go to state 2, return to it as many times as we like, and then go from state 2 to state 0 at the end.

At first sight this does not seem to have simplified matters at all. But we have now gained control over how we use the state 2 because we can use this observation to obtain the following equality.

$$\mathcal{L}_{0 \rightarrow 0} = \mathcal{L}_{0 \rightarrow 0}^{\leq 2} = \mathcal{L}_{0 \rightarrow 0}^{\leq 1} \cup \mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^* \cdot \mathcal{L}_{2 \rightarrow 0}^{\leq 1}$$

We have introduced new notation here, using $\mathcal{L}_{i \rightarrow j}^{\leq 1}$ to mean that we go from state i to state j while using only states 0 and 1 in between, that is, states with a number ≤ 1 . While our expression has grown the individual languages are now easier to calculate.

- $\mathcal{L}_{0 \rightarrow 0}^{\leq 1}$. We cannot go from state 0 to state 0 while only using states 0 and 1 in between, other than just staying in state 0, which corresponds to following the empty word ϵ through the automaton. Hence this language is equal to $\{\epsilon\}$.
- $\mathcal{L}_{0 \rightarrow 2}^{\leq 1}$. To go from state 0 to state 2 without using state 2 in between we must see c followed by b . Hence this language is equal to $\{cb\}$.
- $\mathcal{L}_{2 \rightarrow 2}^{\leq 1}$. To go from state 2 to state 2 without using state 2 in between there are two possibilities:
 - we can go from state 2 to state 1 and back, seeing cb along the way or
 - we can go from state 2 to state 0 to state 1 to state 2 and see acb along the way. (It would be a good idea now to convince yourself that there is no other possibility.)

Hence this language is equal to $\{cb\} \cup \{acb\} = \{cb, acb\}$.

- $\mathcal{L}_{2 \rightarrow 0}^{\leq 1}$. The only way of getting from state 2 to state 0 using only states 0 and 1 in between is the direct route, which means we must see a . Hence this language is equal to $\{a\}$.

Putting all these together we get the following.

$$\begin{aligned} \mathcal{L}_{0 \rightarrow 0} = \mathcal{L}_{0 \rightarrow 0}^{\leq 2} &= \mathcal{L}_{0 \rightarrow 0}^{\leq 1} \cup \mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^* \cdot \mathcal{L}_{2 \rightarrow 0}^{\leq 1} \\ &= \{\epsilon\} \cup \{cb\} \cdot \{cb, acb\}^* \cdot \{a\} \end{aligned}$$

We can now read a pattern⁴ off from this expression, namely

$$\epsilon|cb(cb|acb)^*a,$$

and we have

$$\mathcal{L}_{0 \rightarrow 0} = \mathcal{L}(\epsilon|cb(cb|acb)^*a).$$

Note that this is not the simplest pattern we could possibly have used— $(c(bc)^*ba)^*$ would also have worked. The advantage of the algorithm is that it *always* gives you a pattern in the end, no matter how complicated the automaton.

In this case the languages we were required to calculate were not so difficult to work out. If we had started with a more complicated automaton we would have had to employ the same trick again, namely taking another state and controlling how that is used. The second example shows how this works.

We do yet have to calculate the second language, $\mathcal{L}_{0 \rightarrow 2}$. Again we do so by controlling when we use state 2, but this time it is slightly easier:

To get from state 0 to state 2 we can

- go to state 2 and return to it as many times as we like.⁵

This leads us to the following equality.

$$\mathcal{L}_{0 \rightarrow 2} = \mathcal{L}_{0 \rightarrow 2}^{\leq 2} = \mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^*.$$

Moreover, we already have worked out the languages required here, so we can continue directly to write the following.

$$\begin{aligned} \mathcal{L}_{0 \rightarrow 2} &= \mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^* \\ &= \{cb\} \cdot \{cb, acb\}^* \end{aligned}$$

From this we can read off a regular expression, namely $cb(cb|acb)^*$, and we get

$$\mathcal{L}_{0 \rightarrow 2} = \mathcal{L}(cb(cb|acb)^*).$$

Altogether that means

$$\begin{aligned} \mathcal{L} &= \mathcal{L}_{0 \rightarrow 0} \cup \mathcal{L}_{0 \rightarrow 2} \\ &= \mathcal{L}(\epsilon|cb(cb|acb)^*a) \cup \mathcal{L}(cb(cb|acb)^*) \\ &= \mathcal{L}(\epsilon|cb(cb|acb)^*a|cb(cb|acb)^*) \end{aligned}$$

If we like we can simplify that as follows.

$$\mathcal{L}(\epsilon|cb(cb|acb)^*a|cb(cb|acb)^*) = \mathcal{L}(\epsilon|cb(cb|acb)^*(a|\epsilon))$$

Alternatively, with a bit of experience one might read off this regular expression directly from the automaton:


$$(c(bc)^*ba)^*(\epsilon|c(bc)^*b)$$

Exercise 36. Automata to Regular Expressions

Carry out the algorithm just described for the following automaton. Use the notation with the various languages $\mathcal{L}_{j \rightarrow i}^{\leq k}$ for the first two or three lines at least, even if you can read off a pattern directly. If you want to compare your solution to the model answers you should adopt the numbering given in the picture.


⁴If you find this difficult go back and do Exercises 1, 2, and 8 and have another look at page 18.

⁵It would be a good idea to work out why this is different from the case of language $\mathcal{L}_{0 \rightarrow 0}$.




Algorithm 2, second example

Let's look at a more complicated example.⁶



Again we have to number the states. The most ‘complicated’ state is the one on the top right, and once that is removed there is not much difference, so we use the following numbering.



What is confusing about this automaton is not its size in terms of the number of states, but the way the transitions criss-cross it.

In order to find all the words accepted by the automaton we have to identify all the words that

- when starting in state 0 end up in state 0. We call the resulting language $\mathcal{L}_{0 \rightarrow 0}$.

We also need all the words that

- when starting in state 0 end up in state 2. We call the resulting language $\mathcal{L}_{0 \rightarrow 2}$.

In other words the language \mathcal{L} recognized by the automaton can be calculated as

$$\mathcal{L} = \mathcal{L}_{0 \rightarrow 0} \cup \mathcal{L}_{0 \rightarrow 2}.$$

This now allows us to calculate these two languages separately, which already makes the task a little bit easier. Note that what we have done is to split up the recognized language into all the

⁶Note that the automaton from the first example is a ‘sub-automaton’ of this one—all we have done here is to add some further transitions.

languages of the form $\mathcal{L}_{q_\bullet \rightarrow q}$, where q_\bullet is the start state and q ranges over all the accepting states (that is all the elements of F).

In general we would therefore write for the language \mathcal{L} recognized by some automaton $(Q, q_\bullet, F, \delta)$

$$\mathcal{L} = \bigcup_{q \in F} \mathcal{L}_{q_\bullet \rightarrow q}.$$

But finding all the different paths that lead from 0 to 0, or from 0 to 2 is still pretty tough. The way we simplify that is by taking the state with the highest index, namely 2, out of consideration as follows.

Every path from the state 0 to the state 0 can do one of the following:

- It either doesn't use the state 2 at all or
- it goes from the state 0 to the state 2, then goes back to the state 2 as often as it likes, and ultimately goes to the state 0.

At first sight this doesn't look like a very useful observation. But what we have done now is to break up any path that starts at the state 0 and finishes at the state 0 into a succession of paths that *only use the state 2 at controlled points*.

We use the same notation as before: All words that follow a path that goes from state 0 to state 0 while only using states 0 and 1 (but not state 2) in between make up the language $\mathcal{L}_{0 \rightarrow 0}^{\leq 1}$. This works similarly for other start or end states. Reformulating our last observation means then that every word that follows a path from state 0 to state 0 satisfies one of the following:

- It either is an element of $\mathcal{L}_{0 \rightarrow 0}^{\leq 1}$ or
- it is an element of

$$\mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^* \cdot \mathcal{L}_{2 \rightarrow 0}^{\leq 1}.$$

That means that we have the equality

$$\mathcal{L}_{0 \rightarrow 0} = \mathcal{L}_{0 \rightarrow 0}^{\leq 2} = \mathcal{L}_{0 \rightarrow 0}^{\leq 1} \cup \mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^* \cdot \mathcal{L}_{2 \rightarrow 0}^{\leq 1}.$$

While the equality may appear to make things more confusing at first sight, we now have languages which we can more easily determine on the right hand side of the equality.

We now have the choice between trying to determine the languages on the right directly, or applying the same idea again.

- $\mathcal{L}_{0 \rightarrow 0}^{\leq 1}$. How do we get from state 0 to state 0 only using states 0 and 1? The simple answer is that we can't move there, but we are already there so $\mathcal{L}_{0 \rightarrow 0}^{\leq 1} = \{\epsilon\}$.
- $\mathcal{L}_{0 \rightarrow 2}^{\leq 1}$. Going from state 0 to state 2 using only states 0 and 1 can be done in two ways, either directly using the letter b or *via* state 1 using ab or cb . Hence $\mathcal{L}_{0 \rightarrow 2}^{\leq 1} = \{b, ab, cb\}$.
- $\mathcal{L}_{2 \rightarrow 2}^{\leq 1}$. This is more complicated. Instead of trying to work this out directly we apply our rule again: When going from state 2 to state 2 using only states 0 and 1 we can either go directly from state 2 to state 2 or we can go from state 2 to state 1, return to state 1 as often as we like using only state 0, and then go from state 1 to state 2. In other words we have.

$$\mathcal{L}_{2 \rightarrow 2}^{\leq 1} = \mathcal{L}_{2 \rightarrow 2}^{\leq 0} \cup \mathcal{L}_{2 \rightarrow 1}^{\leq 0} \cdot (\mathcal{L}_{1 \rightarrow 1}^{\leq 0})^* \cdot \mathcal{L}_{1 \rightarrow 2}^{\leq 0}.$$

We now read off $\mathcal{L}_{2 \rightarrow 2}^{\leq 0} = \{b, ab\}$, $\mathcal{L}_{2 \rightarrow 1}^{\leq 0} = \{aa, ac, c\}$, $\mathcal{L}_{1 \rightarrow 1}^{\leq 0} = \{\epsilon\}$ and $\mathcal{L}_{1 \rightarrow 2}^{\leq 0} = \{b\}$. That gives us

$$\begin{aligned} \mathcal{L}_{2 \rightarrow 2}^{\leq 1} &= \{b, ab\} \cup \{aa, ac, c\} \cdot \{\epsilon\}^* \cdot \{b\} \\ &= \{b, ab\} \cup \{aa, ac, c\} \cdot \{\epsilon\} \cdot \{b\} \\ &= \{b, ab, aab, acb, cb\}. \end{aligned}$$

One step in this calculation merits further explanation, namely \emptyset^* , which is calculated to be $\{\epsilon\}$ on page 13.

- $\mathcal{L}_{2 \rightarrow 0}^{\leq 1}$. This one we can read off directly, it is $\{a\}$.

Altogether we get

$$\begin{aligned}\mathcal{L}_{0 \rightarrow 0} &= \{\epsilon\} \cup \{b, ab, cb\} \cdot \{b, ab, aab, acb, cb\}^* \cdot \{a\} \\ &= \mathcal{L}(\epsilon) \cup \mathcal{L}(b|ab|cb) \cdot (\mathcal{L}(b|ab|aab|acb|cb))^* \cdot \mathcal{L}(a) \\ &= \mathcal{L}(\epsilon|(b|ab|cb)(b|ab|aab|acb|cb)^*a).\end{aligned}$$

See the top of page 18 for an explanation of the last step.

This leaves the calculation of $\mathcal{L}_{0 \rightarrow 2}$. We once again apply the earlier trick and observe this time that a path from state 0 to state 2 will have to reach state 2 (for the first time) using only states 0 and 1 along the way, and then it can return to state 2 as often as it likes, giving

$$\mathcal{L}_{0 \rightarrow 2} = \mathcal{L}_{0 \rightarrow 2}^{\leq 2} = \mathcal{L}_{0 \rightarrow 2}^{\leq 1} \cdot (\mathcal{L}_{2 \rightarrow 2}^{\leq 1})^*.$$

Now $\mathcal{L}_{0 \rightarrow 2}^{\leq 1}$ we already calculated above, it is equal to $\{b, ab, cb\}$. We also know already that $\mathcal{L}_{2 \rightarrow 2}^{\leq 1} = \{b, ab, aab, acb, cb\}$. Hence

$$\mathcal{L}_{0 \rightarrow 2} = \{b, ab, cb\} \cdot \{b, ab, aab, acb, cb\}^* = \mathcal{L}((b|ab|cb)(b|ab|aab|acb|cb)^*).$$

Hence the language recognized by the automaton is

$$\begin{aligned}\mathcal{L}_{0 \rightarrow 0} \cup \mathcal{L}_{0 \rightarrow 2} &= \mathcal{L}(\epsilon|(b|ab|cb)(b|ab|aab|acb|cb)^*a) \cup \mathcal{L}((b|ab|cb)(b|ab|aab|acb|cb)^*) \\ &= \mathcal{L}(\epsilon|((b|ab|cb)(b|ab|aab|acb|cb)^*a)|((b|ab|cb)(b|ab|aab|acb|cb)^*)) \\ &= \mathcal{L}(\epsilon(b|ab|cb)(b|ab|aab|acb|cb)^*(\epsilon|a)),\end{aligned}$$

and a regular expression giving the same language is

$$\epsilon|(b|ab|cb)(b|ab|aab|acb|cb)^*(\epsilon|a).$$

Algorithm 2, general case

It is time to generalize what we have done into an algorithm, known as Algorithm 2. We assume that the states in the automaton are given by $\{0, 1, \dots, n\}$, and that j is the start state, where $0 \leq j \leq n$. Then the language \mathcal{L} of all the words accepted by the automaton is equal to

$$\bigcup_{i \in F} \mathcal{L}_{j \rightarrow i} = \bigcup_{i \in F} \mathcal{L}_{j \rightarrow i}^{\leq n}$$

where $\mathcal{L}_{j \rightarrow i}$ is the language of all words that, when starting in state j end in state i – recall that F is the set of accepting states of the automaton. Since a word is accepted if and only if it ends in an accepting state the above equality is precisely what we need.

We can now think of some $\mathcal{L}_{j \rightarrow i}$ as equal to $\mathcal{L}_{j \rightarrow i}^{\leq n}$: the language of all words that, when starting in state j , end up state i is clearly the language of all words that do so when using any of the states in $\{0, 1, \dots, n\}$. In general, $\mathcal{L}_{j \rightarrow i}^{\leq k}$ is the language of all those words that, when starting in state j end in state i , use only states with a number less than or equal to k in between. It is the languages of the form $\mathcal{L}_{j \rightarrow i}^{\leq k}$ for which we can find expressions that reduce k by 1: Any path that goes from state j to state i using only states with numbers at most k will

- either go from state j to state i only using states with number at most $k-1$ in between (that is, not use state k at all)

- or go from state j to state k (using only states with number at most $k - 1$ in between), return to state k an arbitrary number of times, and then go from state k to state i using only states with number at most $k - 1$ in between.

Hence we have

$$\mathcal{L}_{j \rightarrow i}^{\leq k} = \mathcal{L}_{j \rightarrow i}^{\leq k-1} \cup \mathcal{L}_{j \rightarrow k}^{\leq k-1} \cdot (\mathcal{L}_{k \rightarrow k}^{\leq k-1})^* \cdot \mathcal{L}_{k \rightarrow i}^{\leq k-1}.$$

We note that if $j = k$ or $i = k$ then in the above expression only two different languages appear, namely

$$\text{if } j = k \quad \text{then} \quad \mathcal{L}_{j \rightarrow i}^{\leq k-1} = \mathcal{L}_{k \rightarrow i}^{\leq k-1} \quad \text{and} \quad \mathcal{L}_{j \rightarrow k}^{\leq k-1} = \mathcal{L}_{k \rightarrow k}^{\leq k-1}$$

and

$$\text{if } i = k \quad \text{then} \quad \mathcal{L}_{j \rightarrow i}^{\leq k-1} = \mathcal{L}_{j \rightarrow k}^{\leq k-1} \quad \text{and} \quad \mathcal{L}_{k \rightarrow i}^{\leq k-1} = \mathcal{L}_{k \rightarrow k}^{\leq k-1}.$$

Thus we get slightly simpler expressions (compare $\mathcal{L}_{0 \rightarrow 2}$ in the above example):

$$\begin{aligned} \mathcal{L}_{j \rightarrow i}^{\leq j} &= (\mathcal{L}_{j \rightarrow j}^{\leq j-1})^* \mathcal{L}_{j \rightarrow i}^{\leq j-1} \\ \mathcal{L}_{j \rightarrow i}^{\leq i} &= \mathcal{L}_{j \rightarrow i}^{\leq i-1} (\mathcal{L}_{i \rightarrow i}^{\leq i-1})^* \end{aligned}$$

Because the number of states involved in the paths in questions goes down the languages on the right hand side of the equality are easier to calculate. If we have to we can keep applying the above equality until we have reached languages where the only state allowed on the way is a state with number below 0—that is, no state at all. Of course the overall expression we have will become longer and longer, so it pays to read off languages from the automaton as soon as that is possible. Once we have reached paths which are not allowed to use any other states we have the following. For $j = i$:

$$\mathcal{L}_{i \rightarrow i}^{\leq -1} = \begin{cases} \{\epsilon\} & \text{if there is no transition from } i \text{ to } i \\ \{\epsilon, x_1, x_2, \dots, x_l\} & \text{if the } x_i \text{ are all the labels on the transition from } i \text{ to itself.} \end{cases}$$

For $j \neq i$:

$$\mathcal{L}_{j \rightarrow i}^{\leq -1} = \begin{cases} \emptyset & \text{if there is no transition from } j \text{ to } i \\ \{x_1, x_2, \dots, x_l\} & \text{if these are all the labels on the transition from } j \text{ to } i. \end{cases}$$

There are a number of equalities that allow us to simplify the expression we obtain in this way.

- $\emptyset \cup \mathcal{L} = \mathcal{L} = \mathcal{L} \cup \emptyset$;
- $\emptyset^* = \{\epsilon\} = \{\epsilon\}^*$;
- $\emptyset \cdot \mathcal{L} = \emptyset = \mathcal{L} \cdot \emptyset$;
- $\{\epsilon\} \cdot \mathcal{L} = \mathcal{L} = \mathcal{L} \cdot \{\epsilon\}$
- $(\{\epsilon\} \cup \mathcal{L})^* = \mathcal{L}^* = (\mathcal{L} \cup \{\epsilon\})^*$.


Once we have an expression consisting entirely of the simple languages of the form $\mathcal{L}_{j \rightarrow i}^{\leq -1}$ we can insert regular expressions generating these, and use the rules from page 18 to get one regular expression for the overall language.

The advantage of having a general algorithm that solves the problem is that one can now write a program implementing it that will work for all automata, even large ones.


For the marked exercises below, you should use the algorithm as described – even if it’s “obvious” what the pattern might be.

Exercise 37. *Automata to Regular Expressions*


Using Algorithm 2 find a pattern for the following automaton. Can you then also read off a pattern from the automaton? Does it match the one constructed?

**Exercise 38.** *Automata to Regular Expressions*


Repeat Exercise 37 for the following automaton.

**Exercise 39.** *Automata to Regular Expressions*


Repeat Exercise 37 for the following automaton.

**Exercise 40.** *Automata to Regular Expressions*


Repeat Exercise 37 for the following automaton.

**Exercise 41.** *Automata to Regular Expressions*


Give a regular expression defining the language recognized by the following automaton using Algorithm 2.

**Exercise 42.** Automata to Regular Expressions

Repeat Exercise 41 for the following automaton.

**Exercise 43.** Automata to Regular Expressions

Repeat Exercise 41 for the following automaton.



Recall that the way you number the states has an impact on how many steps of the algorithm you will have to apply! Once you have solved the exercises above, try numbering the states differently and redoing the task.

By giving the algorithm we have established the following result.

Theorem 3.2. For every DFA there is a regular expression such that the language recognized by the DFA is the language defined by the regular expression.

We get the following consequence thanks to Theorem 3.1.

Theorem 3.3. For every DFA or NFA it is the case that the language it recognizes is regular.

3.6.1 An Alternative Approach: GNFA's

Given a DFA, Algorithm 2 provides us with a way of deriving a regular expression such that the language defined by the regular expression is the language recognised by the automaton⁷.

⁷Actually, this also works for an NFA

You should not be surprised to hear that there are other ways in which we can derive such a regular expression⁸. The notion of a Generalised Nondeterministic Finite Automaton or GNFA can also be used to do this. We give here a brief description of GNFA's and how they can be used to derive a regular expression for an automaton⁹.

Definition 10. Let Σ be a finite alphabet of symbols and $RegExp(\Sigma)$ the set of regular expressions over that alphabet. A Generalised Nondeterministic Finite Automaton GNFA over Σ consists of the following:

- A finite non-empty set Q of states;
- a particular element of Q called the start state (denoted q_\bullet);
- a particular element of Q called the accepting state (denoted q_a);
- a transition function $\delta : (Q - \{q_a\}) \times (Q - \{q_\bullet\}) \rightarrow RegExp(\Sigma)$.

The key things to note here are that the transitions are labelled with regular expressions, there is a single accepting state, and there is a transition from every state to every other state (apart from the start state and accepting state which have no incoming or outgoing transitions respectively). The notion of acceptance is similar to that for a DFA, but in this case, a word w is accepted if we can find a sequence of states q_0, q_2, \dots, q_n s.t. q_0 is q_\bullet , q_n is q_a , we can split w into a sequence of subwords $w = w_1 w_2 \dots w_n$ and for each w_i , the regular expression that labels the transition from q_{i-1} to q_i matches the word.

There is a straightforward process that translates a DFA to a GNFA that recognises the same language. We add a new start state with an transition labelled ϵ to the existing start state, and a new accepting state with a transition labelled ϵ from the old accepting states. If there are any states with multiple transitions between them labelled x, y, \dots, z , we replace those with a single transition labelled $x|y| \dots |z$. Finally, if there are any states that are not connected with a transition (as required in the definition above), we add a transition labelled with \emptyset . Hopefully you can see that this process produces a GNFA that (using the definition of acceptance introduced above) accepts the same language as the original DFA.

There is also a straightforward process that takes a GNFA and “removes” a state from the machine, producing another GNFA that recognises the same language. By applying this process repeatedly, we end up with a GNFA with only two states (the start state and accepting state), and a single transition from the start state to the accepting state labelled with a regular expression. This regular expression defines the language recognised by the machine (and thus by the original DFA). If you are interested in knowing more, this is covered in some depth in Sipser Chapter 1 (p.70 onwards).

3.7 From patterns to automata

We have a way of going from an automaton to a pattern that we can communicate to a computer, so a natural question is whether one can also go in the opposite direction. This may sound like a theoretical concern at first sight, but it is actually quite useful to be able to derive an automaton from a pattern. That way, if one does come across a pattern that one doesn't entirely understand one can turn it into an automaton. Also, changing existing patterns so that they apply to slightly different tasks can often be easier done by first translating them to an automaton.

For some patterns we can do this quite easily.

⁸Note here that we have referred to *a* regular expression. Is there only one such regular expression for a given automaton?

⁹Note that this section is not part of the examinable material for the unit.

Exercise 44. DFAs

Design DFAs over the alphabet $\{a, b, c\}$ that recognize the languages defined by the following patterns.

- (a) $(a|b)cc$.
- (b) $cc(a|b)$.
- (c) $aa|bb|cc$.
- (d) $c(a|b)^*c$.

Now assume that instead we want to recognize all words that contain a substring matching those patterns. How do you have to change your automata to achieve that?

For slightly more complicated patterns we can still do this without too many problems, provided we are allowed to use NFAs.

Exercise 45. NFAs for Patterns

Design NFAs over the language $\{0, 1\}$ that recognize the languages defined by the following patterns.

- (a) $(00)^*|(01)^*$
- (b) $(010)^*|0(11)^*$


Now turn your NFAs into DFAs.

However, in general it can be quite difficult to read off an automaton from a pattern.¹⁰ We therefore introduce an algorithm that works for all regular expressions. This algorithm is *recursive*, and it builds on Definition 1, making use of the recursive structure of patterns. It is very easy to build automata for the base cases. However, to build automata for the constructors of patterns alternative, concatenation, and star, we need to be a bit cleverer.


Algorithm 3, example

Assume that we want to build an automaton for the regular expression a^*b^* based on already having automata for the patterns a^* and b^* .

The latter are quite easily constructed:



Now we would somehow like to stick these automata together, so that the first part accepts the a^* part of our word, and the second part the b^* part. For that we'd like to have a way of going from the left state above to the right state above in a way that doesn't consume any of the letters of our word, something like this:



¹⁰You can always write down some reasonably complex patterns and give it a go.

That then is precisely what we do: We generalize our notion of automaton to include transitions labelled not by a letter from our alphabet, but by the empty word ϵ .

Definition 11. Let Σ be an alphabet not containing ϵ . An **NFA with ϵ -transitions over Σ** is an NFA over the alphabet Σ that may have transitions labelled with ϵ . Hence the transition relation δ relates pairs of the form (q, x) , where q is a state and x is either an element of Σ or equal to ϵ , to states.

We now have to worry about what it means for an NFA with ϵ -transitions to accept a word. Whenever there is a transition labelled with ϵ in the automaton we are allowed to follow it without matching the next letter in our word.

Definition 12. A word $s = x_1 \cdots x_n$ over Σ is **accepted by the NFA with ϵ -transitions** $(Q, q_\bullet, F, \delta)$ over Σ if there are states

$$q_0 = q_\bullet, q_1, \dots, q_l$$

and, for all $1 \leq i \leq n$, transitions

$$q_{m_{i-1}} \xrightarrow{\epsilon} q_{m_{i-1}+1} \xrightarrow{\epsilon} \dots \xrightarrow{\epsilon} q_{m_i-1} \xrightarrow{x_i} q_{m_i}$$

as well as transitions

$$q_{m_n} \xrightarrow{\epsilon} q_{m_n+1}^1 \xrightarrow{\epsilon} \dots \xrightarrow{\epsilon} q_l,$$

such that $m_0 = 0$ and such that q_l is an accepting state. Here in each case the number of ϵ -transitions may be 0. The **language recognized by an NFA with ϵ -transitions** is the set of all words it accepts.

While we do need NFAs with ϵ -transitions along the way to constructing an automaton from a pattern we do not want to keep the ϵ -transitions around since they make the automata in question much more confusing. We introduce an algorithm here that removes the ϵ -transitions from such an automaton. For that let Σ be an arbitrary alphabet not containing ϵ .

Before we turn to describing the **general case of Algorithm 3** we investigate the algorithm that removes ϵ -transitions.

Algorithm 4, general case

Algorithm 4 is a case where it makes sense to give the general case and then look at an example. Let $(Q, q_\bullet, F, \delta)$ be an NFA with ϵ -transitions over some alphabet Σ . We define an NFA $(Q, q_\bullet, F', \delta')$ without ϵ -transitions over the same alphabet as follows:

- Let

$$F' = F \cup \{q \in Q \mid \exists q = q_0, \dots, q_n \in Q \text{ with } q_i \xrightarrow{\epsilon} q_{i+1} \text{ for all } 0 \leq i \leq n-1 \text{ and } q_n \in F\}.$$

In other words we add to F those states for which we can get to an accepting state by following only ϵ moves.

- For $q, q' \in Q$ and $x \in \Sigma$, the transition relation δ' relates (q, x) to q' if and only if there are states $q = q_1, q_2 \dots q_n = q'$ such that


$$q = q_1 \xrightarrow{\epsilon} q_2 \cdots q_{n-2} \xrightarrow{\epsilon} q_{n-1} \xrightarrow{x} q_n = q'.$$

In other words for $x \in \Sigma$ and $q, q' \in Q$ there is a transition labelled x from q to q' in δ' if it is possible to get from q to q' in δ by following an arbitrary number of ϵ -transitions followed by one transition labelled x .

The resulting automaton may then contain unreachable states that can be safely removed. It recognizes precisely the same language as the automaton we started with.


Algorithm 4, example

Here is an example. Consider the following automaton with ϵ -transitions.



We copy the states as they are, and create some new accepting states as follows:


Pick a non-accepting state. If from there we can reach an accepting state (in the original automaton) *using only ϵ -transitions*, we make this state an accepting state.




For the example, this means the initial state and the third state from the left are now accepting states.




We then copy all the transitions labelled with a letter other than ϵ from the original automaton.




We now look at each of these transitions in turn. Take the transition labelled a . We add a transition labelled a from state i to state j if it is possible to get from state i to state j using a number of ϵ -transitions followed by the transition labelled a .




These give rise to the following transitions for our automaton:




This gives us two new transitions included below.




We now do the same for the transition labelled b .



This gives us one new transition, but we already have a transition from the initial state to the third state from the left, so we can just add the new label b to that.¹¹



Now we may safely remove the unreachable states and any transitions involving them.




We now have an automaton without ϵ -transitions that accepts precisely the same words as the original. What is the language that this automaton accepts? Did you find it easier to read off the language from this automaton or the original?

¹¹We could also make do with just one loop transition being drawn—I only left in two to make the process clearer.

Exercise 46. ϵ Removal

Turn the following automaton into one that does not have ϵ -transitions.



There is further opportunity for practising this algorithm below.

Theorem 3.4. *For every NFA with ϵ -transitions there is an NFA over the same alphabet that recognizes the same language.*

Algorithm 3, general case

We are now ready to look at the process by which we recursively turn a pattern over some alphabet Σ into an NFA with ϵ -transitions over the same alphabet, known as Algorithm 3. Recall, that as discussed on page 17, our regular expressions are defined in a recursive fashion. Thus we can define an operation over a regular expression by providing a definition for the base cases and then how to apply concatenation, alternative and Kleene star.

The pattern \emptyset


An automaton recognizing the language defined by this pattern is given below.



An automaton that accepts no word at all.¹²

The pattern ϵ


An automaton recognizing the language defined by this pattern is given below.



An automaton that accepts precisely the word ϵ .¹³

The pattern x for $x \in \Sigma$

An automaton recognizing the language defined by this pattern is given below.



¹²We could leave out the transition entirely.


¹³Again we could leave out the transition and the right hand state, which is a dump state.

An automaton that accepts precisely the word x .

For the next two cases assume that we have patterns p_1 and p_2 as well as automata A_1 and A_2 such that A_i recognizes the language defined by p_i .

Concatenation


An automaton recognizing the language defined by the pattern p_1p_2 is given below. Here we turn every accepting state of A_1 into a non-accepting state and draw an ϵ -transition from it to the start state of A_2 . The only accepting states in the new automaton are those from A_2 .



In simple cases where the first automaton has a single accepting state, we sometimes omit the ϵ transformation.

Alternative

An automaton accepting the language defined by the pattern $p_1|p_2$ is given below. We add a new start state and connect it with ϵ -transitions to the start states of A_1 and A_2 (so these are no longer start states).



Kleene Star

We assume that we have a pattern p and an automaton A that recognizes the language defined by p .

An automaton accepting the language defined by the pattern p^* is given below. Given an automaton A we introduce a new start state. This state is accepting, and it has an ϵ -transition to the old start state. Further we introduce an ϵ -transition from each accepting state to the old start state.



Be careful when constructing an automaton for a Kleene star expression. A common mistake is to omit the new start state.



Take some time to convince yourself that these new automata are, indeed accepting the languages defined by the resulting patterns.

Exercise 47. NFAs



Refer to your solutions to Exercise 44 to do the following. Design NFAs with ϵ -transitions for the following:

- (a) All words over $\{a, b, c\}$ that match the pattern $(aa|bb|cc)(a|b)cc$.
- (b) All words over $\{a, b, c\}$ that match the pattern $((a|b)cc)|(cc(a|b))$.
- (c) All words over $\{a, b, c\}$ that match the pattern $(aa|bb|cc)^*$.



We illustrate the process with an example. Consider $(a(ba)^*|abaa)^*$. We construct an automaton that accepts the language given by that pattern. For the patterns a and b we have the automata





Hence for ba and $abaa$, respectively, we have




For the left hand automaton we apply the Kleene star, so on the left there's an automaton for $(ba)^*$ now.




In another step we turn the left hand automaton into one for $a(ba)^*$.



We now combine the two automata using the rule for alternative to obtain the following automaton for $(a(ba)^*|abaa)^*$:



It remains to apply the Kleene star to the whole automaton in order to obtain one for $(a(ba)^*|abaa)^*$.



Note that this automaton has more ϵ -transitions than required, and in practice it is fine to leave out ones that are clearly redundant. Note that we have constructed an NFA with ϵ -transitions, so we have to apply Algorithm 4 to remove the ϵ -transitions, and then Algorithm 1 to obtain a DFA. In practice strictly following Algorithm 3 from page 3.7 introduces rather more ϵ -transitions than strictly necessary. In particular when concatenating very simple automata, the ϵ -transition connecting them can often be safely left out. I am happy for you to do so in assessed work.

Exercise 48. *Pattern to NFA+ ϵ*

Use Algorithm 3 described on page 54 to construct an NFA with ϵ -transitions that recognizes the language defined by the pattern $(a|b)^*a(a|b)$.

Exercise 49. *ϵ Removal*

Take your result from Exercise 48 and remove the ϵ -transitions by applying Algorithm 4 on page 51. Now remove all unreachable states.

Exercise 50. *NFA to DFA*

Take your result from Exercise 49 and turn it into a DFA by using Algorithm 1 on page 36. How many states would your DFA have if you didn't remove unreachable states first? How many states would your DFA have if you didn't only draw the reachable ones?

Exercise 51. *DFA's for Patterns*

Construct DFA's that recognize the languages defined by the following patterns. You may either just write down automata, or use Algorithms 3, 4 and 1 to do this.

- (a) $(ab|a)^*$
- (b) $(01)^*|001^*$

Exercise 52. *DFA's for Patterns*

Repeat Exercise 51 for the following patterns.

- (a) $((00)^*11|01)^*$
- (b) $((ab)^*b|ab^*)^*$

We have given an idea of how to prove the following.

Theorem 3.5. *For every regular language we can find the following recognizing it:*

- *deterministic finite automaton;*
- *non-deterministic finite automaton;*
- *non-deterministic finite automaton with ϵ -transitions.*

Altogether, we have shown in various parts of this section the following result.

Theorem 3.6. *The following have precisely the same power when it comes to describing languages:*

- *regular expressions,*
- *deterministic finite automata,*
- *non-deterministic finite automata.*

Hence for every regular language we can find a description using any one of the above, and given one description we can turn it into any of the others.

Why is it useful to have several different ways of representing exactly the same problem? Why don't we just stick with regular expressions? The answer is that, although the different mechanisms allow us to describe the same set of languages, they have different characteristics that mean they are suitable for different purposes.

For example, how would you go about implementing a system that recognises regular expressions? You would certainly need some kind of data structure that represents the regular expression. That would require some thought, but should be relatively straightforward. What is harder is working out exactly how to write a *program* or clear sequence of instructions that tell us what to *do* when trying to match a word to a regular expression.

If we consider finite automata, and in particular *deterministic* finite automata, this is a different proposition. We would have to define suitable data structures to represent the states and transitions, but it is much easier to describe *in code* what needs to happen when trying to judge whether or not a word is accepted. The DFA is mechanical, and the transition function effectively tells us what to do next.

A trade off here, of course, is that it is often easier for us to write down regular expressions than it is to describe an equivalent automaton. So as a *user*, I might be happier working with regular expressions, while as an *implementer* I may be happier with automata. But this is ok – as we have seen above, there are well defined algorithms that will convert regular expressions to automata and vice versa. In fact this is exactly what some regular expression matching engines will do – convert a regular expression to a finite state machine which can then be run on the input.

3.8 Properties of regular languages

In Section 2 there are examples of how we can build new regular languages from existing ones. At first sight these may seem like theoretical results without much practical use. However, what they allow us to do is to build up quite complicated languages from simpler ones. This also means that we can build the corresponding regular expressions or automata from simple ones, following established algorithms. That makes finding suitable patterns of automata less error-prone, which can be very important.

If our language is finite to start with then finding a pattern for it is very easy.

Exercise 53. *Finite Languages*

Show that every finite language is regular.

Assume we have two regular languages, \mathcal{L}_1 and \mathcal{L}_2 . We look at languages we can build from these, and show that these are regular.

Concatenation

In Section 2 it is shown that the concatenation of two regular languages is regular. If \mathcal{L}_1 and \mathcal{L}_2 are two regular languages, and p_1 is a regular expression defining the first, and p_2 works in the

same way for the second, then

$$\mathcal{L}_1 \cdot \mathcal{L}_2 = \mathcal{L}(p_1) \cdot \mathcal{L}(p_2) = \mathcal{L}(p_1 p_2),$$

so this is also a regular language. There is a description how to construct an automaton for $\mathcal{L}_1 \cdot \mathcal{L}_2$ from those for \mathcal{L}_1 and \mathcal{L}_2 on page 54.

Kleene star

Again this is something we have already considered. If \mathcal{L} is a regular language then so is \mathcal{L}^* . If p is a regular expression for \mathcal{L} then p^* is one for \mathcal{L}^* , and again our algorithm for turning patterns into automata shows us how to turn an automaton for \mathcal{L} into one for \mathcal{L}^* .

Reversal

If s is a word over some alphabet Σ then we can construct another word over the same alphabet by reading s backwards, or, in other words, reversing it.

Definition 13. For a language \mathcal{L} , we define:

$$\mathcal{L}^R = \{x_n x_{n-1} \cdots x_2 x_1 \mid x_1 x_2 \cdots x_{n-1} x_n \in \mathcal{L}\}.$$

Exercise 54. Reversing Strings

Here are some exercises concerned with reversing strings that offer a good opportunity for practising what you have learned so far. Parts (a) and (e) require material that is discussed in Appendix A, so it only makes sense to do these if you are working through this part of the notes.

- (a) Define the reversal of a string as a recursive function.
- (b) Look at an automaton for the language of all non-empty words over the alphabet $\{a, b\}$ which start with a . How can it be turned into one for the language of all words which end with a ? *Hint: How could we have a given word take the reverse path through the automaton than it would do ordinarily?*
- (c) Look at the language given in Exercise 13 (c). What do the words in its reversal look like? Now look at an automaton for the given language and turn it into one for the reversed language. *Hint: See above.*
- (d) In general, describe informally how, given an automaton for a language \mathcal{L} , one can draw one for the language \mathcal{L}^R .
- (e) Take the formal description of a DFA recognizing a language \mathcal{L} as in Definition 5 and turn that into a formal definition for an NFA with ϵ -transitions which recognizes the language \mathcal{L}^R .

Unions

If we have regular expressions for \mathcal{L}_1 and \mathcal{L}_2 , say p_1 and p_2 respectively, it is easy to build a regular expression for $\mathcal{L}_1 \cup \mathcal{L}_2$, namely $p_1 | p_2$. But how do we build an automaton for $\mathcal{L}_1 \cup \mathcal{L}_2$ from those for \mathcal{L}_1 and \mathcal{L}_2 ? We have already seen how to do that as well—form an NFA with a new start state which has ϵ -transitions to the (old) start states of the two automata, as illustrated on page 54. If we like we can then turn this NFA with ϵ -transitions into a DFA.

Intersections

This isn't so easy. It's not at all clear how one would go about it using patterns, whereas with automata one can see how it might work. The problem is that we can't say 'first get through the automaton for \mathcal{L}_1 , then through that for \mathcal{L}_2 ': When we have followed the word through the first automaton it has been consumed, because we forget about the letters once we have followed a transition for them. So somehow we have to find a way to let the word follow a path through *both automata at the same time*.

Let's try this with an example: Assume we want to describe the language \mathcal{L} of all words that have an even number of *as* and an odd number of *bs*. Clearly

$$\begin{aligned}\mathcal{L} = & \{s \in \{a,b\}^* \mid s \text{ has an even number of } as\} \\ & \cap \{s \in \{a,b\}^* \mid s \text{ has an odd number of } bs\}.\end{aligned}$$


DFAs for those two languages are easy to construct.



So how do we get an automaton that checks for both properties at the same time? Well, we think of the first automaton as telling us whether the number of *as* is even so far, or odd, and similarly for the second automaton, and the number of *bs*. What we need to do in the automaton we want to construct then is to keep track of the *as* and *bs* at the same time. So we need four states:

$$(as \text{ even}, bs \text{ even}), (as \text{ even}, bs \text{ odd}), (as \text{ odd}, bs \text{ even}), (as \text{ odd}, bs \text{ odd}).$$

The automaton we want looks as follows:



What have we done? We have formed *pairs* of states, a state from the first automaton and a state from the second one. We have then added transitions that literally do follow both automata at the same time.

In general, given DFAs $(Q_1, q_\bullet^1, F_1, \delta_1)$ and $(Q_2, q_\bullet^2, F_2, \delta_2)$ we can form an automaton that recognizes the intersection of the languages recognized by the two DFAs as follows.

- States: $Q_1 \times Q_2$.
- Start state: $(q_\bullet^1, q_\bullet^2)$.
- Accepting states: $F_1 \times F_2$.
- Transition function: δ maps (q_1, q_2) and x to $(\delta_1(q_1, x), \delta_2(q_2, x))$. In other words, there is a transition

$$(q_1, q_2) \xrightarrow{x} (q'_1, q'_2)$$

if and only if there are transitions

$$q_1 \xrightarrow{x} q'_1 \quad \text{and} \quad q_2 \xrightarrow{x} q'_2.$$

We call the result the **product of the two automata**. If we do need a pattern for the intersection we can now apply Algorithm 2 from page 45.

This again illustrates how it can be useful to have different (but equivalent) ways of describing the same thing. The definition of the product automaton is relatively straightforward. Defining intersection or product on regular expressions is not. We can prove the result – the fact that the intersection of two regular languages is regular – using automata and we then know that given two regular expressions there *is* a regular expression that defines the intersection. As we saw earlier, if we wanted to show that the union of two regular languages is regular, we can do this simply by using the $|$ operator.

Exercise 55. *Language Intersection*

Use this construction to draw DFAs recognizing the following languages over the alphabet $\{a, b\}$. First identify the two languages L_1 and L_2 whose intersection you want to form, then draw the automata for those languages, then draw one for their intersection.

- (a) L_1 : all non-empty words that begin with a . L_2 : all non-empty words that end with b .
- (b) L_1 : all words that contain at least two as . L_2 : all words that contain at most one b .

Exercise 56. *Language Intersection*

Repeat Exercise 55 for the languages below.

- (a) L_1 : all words that have even length. L_2 : all words containing an even number of as .
- (b) L_1 : all words that have length at least 3. L_2 : all words whose third symbol (if it exists) is a .

Exercise 57. *Language Intersection*

Repeat Exercise 55 for the languages below.

- (a) L_1 : all words that have odd length. L_2 : all words that contain at least two as .
- (b) L_1 : all words that start with an a . L_2 : all words that do *not* contain the string ba .

Complements

If \mathcal{L} is a language over the alphabet Σ then we may want to consider its complement, that is

$$\Sigma^* - \mathcal{L},$$

the set of all words over Σ that do *not* belong to \mathcal{L} .


Exercise 58. *Language Complement*

This is a question concerned with the complement of a language.

- (a) Consider the language discussed on page 23 of all words which have a 0 in every position that is a multiple of 3. Begin by defining the complement of this language in $\{0, 1\}^*$, either using set-theoretic notation or in English.

Now take the DFA for the original language given on page 24 and *turn it into one* for the complement.

- (b) Consider the following DFA that recognizes the language of all words over $\{0, 1\}$ of length at least 2 whose first letter is a 1 and whose second letter is a 0.



Again, describe the complement of this language either using sets or in English. Then *turn this DFA into one* that recognizes this complement.

Exercise 59. Language Complement

This is a question concerned with the complement of a language.

- (a) Assume you are given a picture that describes a DFA for some arbitrary regular language. Describe in general how you can turn this into a picture for a DFA that recognizes the complement of that language. How does your answer change if you start from an NFA? *Hint: If you find this difficult, then look back at your solutions to 58. How did you have to change the original DFA?*
- (b) Now assume you are given a DFA in the formal description as used in Definition 5, $(Q, q_\bullet, F, \delta)$. Describe how to turn this into a formal description of a DFA recognizing the complement of the original language. (Again you may want to consult Appendix A for this part.)

Reversal via Inductive Proof

Exercise 54 considered the reversal of a language and showed informally that an automaton for a language \mathcal{L} can be turned into one for the language \mathcal{L}^R , and thus that \mathcal{L}^R is regular if \mathcal{L} is regular. As an illustration of the use of inductive proof, we can also demonstrate this through regular expressions. Note that we go through this process in quite a lot of fine detail.

Recall Definition 3.8. For a language \mathcal{L} over Σ we define:

$$\mathcal{L}^R = \{x_n x_{n-1} \cdots x_2 x_1 \mid x_1 x_2 \cdots x_{n-1} x_n \in \mathcal{L}\}.$$

Proposition 3.7. *If \mathcal{L} is regular, then so is \mathcal{L}^R .*

If \mathcal{L} is regular, then we know there is a regular expression p such that $\mathcal{L} = \mathcal{L}(p)$, i.e. every word in \mathcal{L} is matched by the regular expression p . This suggests that a strategy to prove this proposition is to consider reversing that regular expression. But what exactly does it mean to *reverse* a regular expression? We need to be careful – we can't just simply reverse all the symbols in the expression. Take the regular expression $(ab)^*$. A naïve reversal of this would give us $^*(ba)$ which isn't actually a *valid* regular expression.

We first introduce a definition that allows us to talk about reversing words.

Definition 14. *Let $w = x_1 x_2 \cdots x_{n-1} x_n$ be a word over Σ . We define the reversal of w as $w^R = x_n x_{n-1} \cdots x_2 x_1$.*

We note that $\epsilon^R = \epsilon$ and $x^R = x$ for any $x \in \Sigma$. We also note that $w^{RR} = w$.

To define the reversal of a regular expression, we return to Definition 1, which defines regular expressions in a recursive fashion. We use the clauses in that definition to define what it means to reverse each of the base cases (empty pattern, empty word and single letter), and then the cases where new patterns are constructed from existing patterns.

Definition 15. Let p be a regular expression over Σ . The reversal of p , denoted p^R is given as follows.

Empty pattern $\emptyset^R = \emptyset$;

Empty word $\epsilon^R = \epsilon$;

Letters for a letter $x \in \Sigma$, $x^R = x$;

Concatenation if p_1 and p_2 are patterns then $(p_1 p_2)^R = (p_2^R p_1^R)$;

Alternative if p_1 and p_2 are patterns then $(p_1 | p_2)^R = (p_1^R | p_2^R)$;

Kleene star if p is a pattern then $(p^*)^R = (p^R)^*$.

If we follow through the clauses in this definition, we can see that $(ab)^{*R}$ is in fact $(ba)^*$.

Now that we have formally defined reversal applied to words and regular expressions or patterns, we can state the following proposition:

Proposition 3.8. Let p be a pattern over Σ .

If w is a word over Σ such that w matches p , then w^R matches p^R .

We prove this proposition using induction over the definition of regular expressions. As discussed in COMP11120, we have several base and step cases as there are a several primitive terms and ways of putting them together. The principle is the same, however – we first consider base cases, then for our step cases, we make assumptions about the required property holding for a sub-expression.

Base Case: Empty pattern If the pattern p is \emptyset , then there are no words that match the pattern, and the proposition is trivially true;

Base Case: Empty word If the pattern p is ϵ , and w matches p , then w must be the empty word ϵ . Following Definition 15, p^R is ϵ . We have $w^R = \epsilon^R = \epsilon$, which matches ϵ and the proposition holds.

Base Case: Letters If the pattern p is a single $x \in \Sigma$, then p^R is also x . If w matches x , then w must be the word consisting of a single letter x . Then $w^R = x^R = x$ which matches p^R and the proposition holds.

Step Case: Concatenation If the pattern p is a concatenation $(p_1 p_2)$, then our *induction hypothesis* is that the proposition holds for both p_1 and p_2 .

Now, consider a word $w = x_1 x_2 \cdots x_{n-1} x_n$ that matches p . We are required to show that w^R matches p^R . By Definition 2, we know that there must be words s_1 and s_2 such that w is the concatenation of s_1 and s_2 , and that s_1 matches p_1 and s_2 matches p_2 . Our induction hypothesis then tells us that s_1^R matches p_1^R and s_2^R matches p_2^R . Again, consulting the definition for matching a pattern, we can deduce that the concatenation $s_2^R s_1^R$ matches $(p_2^R p_1^R)$. By Definition 15, $(p_2^R p_1^R)^R$ is p^R , so it only remains to show that $s_2^R s_1^R$ is w^R .

We have that w is the concatenation of s_1 and s_2 . Expressing this in a slightly different way, there must be some $0 \leq i \leq n^{14}$ such that

$$s_1 = x_1 x_2 \cdots x_{i-1} x_i \quad \text{and} \quad s_2 = x_{i+1} x_{i+2} \cdots x_{n-1} x_n$$

Then

$$s_2^R = x_n x_{n-1} \cdots x_{i+2} x_{i+1} \quad \text{and} \quad s_1^R = x_i x_{i-1} \cdots x_2 x_1$$

¹⁴Note that the cases where i is 0 or n cover the situations where one of the strings is empty.

so

$$s_2^R s_1^R = x_n x_{n-1} \cdots x_{i+2} x_{i+1} x_i x_{i-1} \cdots x_2 x_1 = w^R$$

as required;

Step Case: Alternative If the pattern p is an alternative $(p_1|p_2)$ then our induction hypotheses is that the proposition holds for both p_1 and p_2 .

Now, consider a word w that matches p . We are required to show that w^R matches p^R . By Definition 2, we know that either 1/ w matches p_1 or 2/ w matches p_2 . If it is the first case, then our induction hypothesis tells us that w^R matches p_1^R . Using Definitions 2 and 15, we see that w^R then matches $(p_1^R|p_2^R)$, which is p^R . Similar reasoning can be applied to case 2, and our proposition holds;

Step Case: Kleene star If the pattern p is an application of the Kleene star $(p_1)^*$, then our induction hypothesis is that the proposition holds for p_1 .

Now, consider a word $w = x_1 x_2 \cdots x_{n-1} x_n$ that matches p . We are required to show that w^R matches $p^R = (p_1^R)^*$. By Definition 2, w can be written as a finite concatenation of words

$$w = s_1 s_2 \cdots s_n$$

such that all words s_i match p_1 . By our induction hypotheses, all s_i^R match p_1^R . Again, by Definition 2, we can construct a word $s_n^R s_{n-1}^R \cdots s_2^R s_1^R$ which then matches $(p_1^R)^*$. It remains to show that $s_n^R s_{n-1}^R \cdots s_2^R s_1^R = w^R$, but this is easily done using similar reasoning to that in the concatenation case above, and our proposition holds.

It is also useful to consider what happens when we apply our reversing process twice to a regular expression.

Proposition 3.9. *Let p be a pattern over Σ . $p^{RR} = p$, i.e. applying the reverse operation twice yields the original regular expression.*

Again, we can approach this via induction over the structure of the regular expressions.

Base Case: Empty pattern If the pattern p is \emptyset , then $\emptyset^{RR} = \emptyset^R = \emptyset$;

Base Case: Empty word If the pattern p is ϵ , then $\epsilon^{RR} = \epsilon^R = \epsilon$;

Base Case: Letters If the pattern p is a single $x \in \Sigma$, then $x^{RR} = x^R = x$;

Step Case: Concatenation If the pattern p is a concatenation $(p_1 p_2)$, then our *induction hypothesis* is that the proposition holds for both p_1 and p_2 . Using the definitions above,

$$(p_1 p_2)^{RR} = (p_2^R p_1^R)^R = (p_1^{RR} p_2^{RR}) = (p_1 p_2)$$

Step Case: Alternative If the pattern p is an alternative $(p_1|p_2)$ then our induction hypotheses is that the proposition holds for both p_1 and p_2 . Then

$$(p_1|p_2)^{RR} = (p_1^R|p_2^R)^R = (p_1^{RR}|p_2^{RR}) = (p_1|p_2)$$

Step Case: Kleene star If the pattern p is an application of the Kleene star, $(p_1)^*$, then our induction hypothesis is that the proposition holds for p_1 .

Then

$$(p_1^*)^{RR} = (p_1^R)^R = (p_1^{RR})^* = (p_1)^*$$

We now return to Proposition 3.7, and consider a regular language \mathcal{L} . By Definition 4, there must be a regular expression p such that $\mathcal{L} = \mathcal{L}(p)$, i.e. \mathcal{L} is the set of all words that match p . We claim that $\mathcal{L}^R = \mathcal{L}(p^R)$, i.e. that it is the set of all words matching p^R . This has two parts to it:¹⁵ 1/ any word in \mathcal{L}^R matches p^R ; and 2/ any word that matches p^R is in \mathcal{L}^R .

Consider a word $w \in \mathcal{L}^R$. Following Definition 3.8, there must be a word $s \in \mathcal{L}$ such that $s^R = w$. We know that s must match p (as $s \in \mathcal{L}$), therefore by Proposition 3.8, $w = s^R$ must match p^R . This gives us our first part of the proof.

Now consider a word w that matches p^R . Proposition 3.8 tells us that w^R will match p^{RR} , which is equal to p by Proposition 3.9. w^R is thus in \mathcal{L} . So w^{RR} is in \mathcal{L}^R by Definition 3.8 and as $w^{RR} = w$ as observed above, w is in \mathcal{L}^R as required, and our result is shown.

3.9 Equivalence of Automata

Here is an issue we have not thought about so far. Is there a way of telling whether two automata recognize the same language? You might particularly care about this if you and a friend have created different automata for the same problem and you're not sure whether you're both correct. We say that two automata are **equivalent** if they recognize the same language. Note that if we have an NFA, and we carry out Algorithm 1 for it, then the resulting DFA is equivalent to the NFA we started with.

If you're lucky, this is easy to decide: If you have both drawn a picture of the 'same' automaton, but have merely put different names in your states, then the automata are what we call *isomorphic*. It is easy to show that isomorphic automata recognize the same language.

There are several ways of addressing the more general issue.

Via complementation and intersection

We actually have a method which allows us to calculate whether two automata define the same language. Assume we have automata A and A' , which recognize the languages \mathcal{L} and \mathcal{L}' respectively. We are interested in the question whether $\mathcal{L} = \mathcal{L}'$. Note that

$$\begin{aligned} \mathcal{L} = \mathcal{L}' & \quad \text{if and only if} \quad \mathcal{L} \subseteq \mathcal{L}' \text{ and } \mathcal{L}' \subseteq \mathcal{L} \\ & \quad \text{if and only if} \quad \mathcal{L} \cap (\Sigma^* - \mathcal{L}') = \emptyset \text{ and } \mathcal{L}' \cap (\Sigma^* - \mathcal{L}) = \emptyset. \end{aligned}$$

But we know how to construct automata for the complement of a language, and we also know how to construct automata for the intersection of two languages, so we can reduce our question to that whether the two automata corresponding to the two languages appearing in the above accept no words at all.

However, note that this is quite a bit of work: We have to

- construct two automata for complements—this is not too much work and does not lead to larger automata and
- we have to construct two intersections—the resulting automata will have a size equal to the product of the sizes of the automata involved.

Hence this procedure is not very practical.

Via simulations

Note that both the methods for deciding equivalence of automata we have mentioned so far work for DFAs, but not for their non-deterministic relations. Clearly the question of whether two NFA recognize the same language is even harder than that for two DFAs. We study a different method for comparing automata.

¹⁵Why do we have to show both of these?

The idea is this: Assume we have two NFAs, say $A = (Q, q_\bullet, F, \delta)$ and $B = (P, p_\bullet, E, \gamma)$. If we can show that for each state q of A there is an ‘analogous’ state p of B , then every word accepted by the first automaton should be accepted by the second one, assuming we get our definition of ‘analogous’ right.

Definition 16. We say that a relation \sim from Q to P is a **simulation** between automata A and B if and only if

- $q_\bullet \sim p_\bullet$;
- if $q \sim p$ for some $q \in Q$ and $p \in P$ then
 - if $q \xrightarrow{x} q'$ then there exists $p' \in P$ such that $p \xrightarrow{x} p'$ and $q' \sim p'$,
 - $q \in F$ implies $p \in E$.

Why are simulations useful? It is because of the following result.

Proposition 3.10. If there is a simulation from one NFA to another then every word accepted by the first is accepted by the second. Therefore in this case we have that the language accepted by the first automaton is a subset of the language accepted by the second one.

Why does this hold? A word $x_0x_1 \cdots x_n$ is accepted by automaton $A = (Q, q_\bullet, F, \delta)$ if and only if there are states

$$q_0 = q_\bullet, q_1, \dots, q_n$$

such that for all $0 \leq i < n$, $q_i \xrightarrow{x_i} q_{i+1}$ and such that $q_n \in F$, that is, q_n is an accepting state. Because there is a simulation \sim from A to $B = (P, p_\bullet, E, \gamma)$ we have that $q_\bullet \sim p_\bullet = p_0$ and because $q_0 \xrightarrow{x_0} q_1$ we may find a state p_1 such that $p_0 \xrightarrow{x_0} p_1$ and $q_1 \sim p_1$ and so that if q_1 is accepting then so is p_1 . So we may construct, one after the other, states p_0, p_1 up to p_n such that for every p_i it is the case that

- $p_i \xrightarrow{x_i} p_{i+1}$,
- $q_i \sim p_i$ and
- q_i accepting implies p_i accepting.

Since q_n is accepting this is also the case for p_n and so the given word is accepted by B . For deterministic automata we have an even stronger result:


Proposition 3.11. Let A and B be DFAs. If the language defined by A is a subset of the language defined by B then there exists a simulation from A to B .

Hence overall we have two theorems, one for NFAs and one for DFAs.


Theorem 3.12. If A and B are NFAs, and there is a simulation from A to B , as well as one in the opposite direction, then A and B accept the same language.

Theorem 3.13. Let A and B be two DFAs. Then A and B accept the same language if and only if there are simulations between them in both directions.


How do we construct simulations in practice?¹⁶ Consider the following automata, assuming we want to construct a simulation from the one on the left to the one on the right.



The two initial states have to be in the relation, which we picture as follows:




Now the first automaton can make a move $00 \xrightarrow{a} 10$, and we must add find a ‘partner’ for 10 in the other automaton to match this transition. The only choice is to connect 10 with 0, pictorially:




The first automaton can make a move $00 \xrightarrow{b} 01$, and to match that we have to ‘partner’ 01 with 1. The state 01 is accepting, but so is 1, so we have satisfied that criterion. We get the following.

¹⁶Note that if you are working from automata that have been drawn, and your source automaton has dump states drawn then you may have to draw a dump state for the target automaton to draw a simulation.




Similarly, we have to find a state matching 11, and that has to be a state which can be reached with b from 10's partner 0, and with a from 01's partner 1, making 1 the only possible choice.




We still have to check that every transition in the automaton on the left can be matched by one between suitable 'partner states' on the right, but this is straightforward.¹⁷

Now assume we would like to construct a simulation in the *opposite direction*, that is, from right to left. The two start states have to be connected.




To match transitions from 0 we have to partner the state 0 with 10 and we have to partner the state 1 with 01. The state 1 is an accepting state, but so is the state 01, hence so far we are on track for constructing a simulation.


¹⁷Note that a state on the left may have more than one partner state on the right—it's coincidence that we do not need more in this example.



But now we have to match the transition from 1 with a to 1, which would require partnering state 1 with state 11. But 1 is an accepting state and may only be partnered with another accepting state. We had no choice in this construction, and since we cannot complete it to give a simulation, there *does not exist* a simulation from the automaton on the right to the one on the left. Here is another example. Consider the following two automata.



We give a simulation from the one on the left to the one on the right:



This can be represented as the following set of pairs:


$$(0, A), (1, B), (2, C), (3, B), (4, A), (5, A), (6, D)$$

If we invert the simulation so that it goes from the automaton on the right to that on the left then in this case we obtain another simulation. Hence the two automata accept the same language. We call a relation which works in both ways a *bisimulation*.

Note that if we have deterministic automata as the target then every simulation may be constructed in this way, one step at a time. One stops when all the states in the source automaton have been matched with a partner in such a way that all transitions in the source automaton can be matched, or when it is impossible to do so. In the first case one has constructed the *unique minimal simulation* from the first automaton to the second, and in the second case one has illustrated that such a simulation does not exist.


Exercise 60. *Simulation*

For the following pair of automata, find simulations going both ways, or argue that one, or both of these, cannot exist.



Exercise 61. *Simulation*


Repeat Exercise 60 for the following pair of automata.



Exercise 62. *Simulation*

For the following pair of automata, find simulations going both ways, or argue that one, or both of these, cannot exist.


Do the automata accept the same language?



Exercise 63. *Simulation*

For the following pair of automata, find simulations going both ways, or argue that one, or both of these, cannot exist.

Do the automata accept the same language?



The notion of simulation allows us to identify when automata are *equivalent* in the sense of recognising the same language. What about regular expressions? Is there a way of determining whether two regular expressions will match exactly the same words?

Of course the answer to this questions is that there is a straightforward mechanism for determining equivalence of regular expressions and it should hopefully be clear to you what that process is.

Exercise 64. *Equivalence of Regular Expressions*

Given two regular expressions p_1 and p_2 describe how you would show that the two expressions are (or are not) equivalent, where p_1 is equivalent to p_2 precisely when they match the same set of words or define the same language.


Via minimization

Consider the two automata below over the alphabet $\{a, b\}$.



Hopefully, you can see that these two automata are equivalent (if this is not obvious to you, try and construct the appropriate bi-simulation). The second has fewer states though, so for some purposes – implementation for example – it may be more desirable.

To further illustrate the point, consider the automaton below.



Again, this automaton accepts the same language, but has even more states. We can go on doing this and can provide an arbitrarily large automaton that is equivalent to our original.

This leads us to consider whether it is possible to construct a *minimal* DFA, where the number of states is as small as possible.

For DFAs, we have the following result:

Theorem 3.14. *For every deterministic automaton A we can find a minimal automaton M such that:*

- M is equivalent to A (recognises the same language).
- For any automaton B that is equivalent to A , the number of states in B will be greater than or equal to the number of states in M .

In addition, if we disregard any naming or labelling of states, M will be unique (it will be *isomorphic* to any other minimal DFA).

This gives us a further way of determining equivalence – two DFAs are equivalent if their minimised automata are isomorphic.

There are several algorithms for calculating the minimal automaton – we will consider one known as the table filling algorithm (for reasons which will become apparent). As an aside, the fact that there are several algorithms should lead us again to more questions. Which of these algorithms is better? And what might better mean in this context? Some aspects of this question will be considered in Part II of the course, when we look at complexity results.

First of all, we define what it means for pairs of states in an automaton to be *distinguishable*.

Definition 17. *Let $A = (Q, q_\bullet, F, \delta)$ be a DFA over an alphabet Σ . For any two states $p, q \in Q$, we say the states are distinguishable if:*


- $p \in F$ and $q \notin F$ or vice versa; or
- for some $x \in \Sigma$, $\delta(p, x)$ is distinguishable from $\delta(q, x)$.

Note that as seen earlier, this is a recursive definition. To put this in simple terms, two states are distinguishable if we can find a word that will lead us to an acceptance from one of the states and non-acceptance from the other. We say that two states are *non-distinguishable* if they are not distinguishable – i.e. there is no word that can result in a different outcome when starting in those states. Note that given this definition, we have that any state p is non-distinguishable from itself¹⁸.

To illustrate this, consider our initial example. The first automaton has three states, $\{0, 1, 2\}$. If we follow our definition, we can see that state 0 is distinguishable from state 1, as 0 is non-accepting, while 1 is accepting. Similarly, states 0 and 2 are distinguishable. What about 1 and 2 though? They are both accepting, so the first clause in Definition 17 does not hold. If we start in state 1, *any* word starting with a will end up in state 2. Also, any word starting in b will end up in state 2. Similarly, if we start in state 2, any word starting with a or b will end up in state 2. The second clause then doesn't hold, so the states 1 and 2 are thus non-distinguishable. In fact, for the examples given at the beginning of this section, all of the accepting states are indistinguishable from each other.

Now, if two states are *non-distinguishable*, i.e. we can't find a word that would result in a different outcome when starting in each state, then we can produce a new automaton with those states merged which will accept the same words.

In our example above, this would mean merging states 1 and 2 to give a new state (let's call it 3), giving us:




¹⁸If you don't see why this must be the case, think about whether it is possible to find a word that results in two different paths through the automaton. Remember this is for a DFA!

which is exactly the second automaton (with a relabelling of the states).

If our automaton was a little bit more complicated, it might be that the identification of two non-distinguishable states allows us to identify further pairs. There is an algorithm called the table filling algorithm that enables us to do precisely this. In simple terms, we loop round, each time checking whether any new pairs have been added, and if so, whether that new information allows us to infer any more pairs. Once we reach a stable state where no new pairs can be found, we stop.

Table filling algorithm: Example

We illustrate the algorithm with an example, using the following automaton:



First, we construct a table that has an entry for each (non-equal) pair of states. So it's actually just under half a table:

0				
1				
2				
3				
	0	1	2	3

For any pairs of states where one is accepting and one is not accepting, we mark the states as distinguished. This means that states 0 and 2 are distinguished from 1 and 3.

0				
1	×			
2		×		
3	×		×	
	0	1	2	3

Now we iterate around our loop. For any pairs of states p and q that are not yet marked as distinguished, we check to see if there is a character x such that there is an edge from p to a state p_x and an edge from q to a state q_x , where p_x and q_x are marked as distinguished. If this is the case, we mark p and q as distinguished (recall Definition 17) and make a note of the fact that we have added some new information.

If we have looked at all relevant pairs and have not added any more information during the loop, then we are finished.


In our example, we need to check the pairs (1,3) and (0,2) as they are the only ones unmarked so far. State 1 only has an edge labelled a that goes to 3, while 3 also has a single outgoing edge labelled a going to 3. So there is nothing to do for this pair. Then, we note that 0 has a transition labelled b that goes to 2, and 2 has a transition labelled b that goes to 1, and 1 and 2 are distinguished. Thus we can mark (0,2) as a distinguished pair. This gives us the following situation:

0				
1	×			
2	×	×		
3	×		×	
	0	1	2	3

We note that we added information during that last pass, so we are not finished yet, as that additional information might have changed the position for pairs of nodes we have already looked at. So we must go around again.

This time there is a single pair to consider – $(1, 3)$. The position with these states remains the same – we cannot find edges that go to states that are distinguishable. This time, there was no information added when looping around, so we are finished (it should be clear that no further changes can occur).

We can now simply read off from the table those pairs of states that are not flagged as distinguishable. Our minimal automaton is then formed by merging those states – 1 and 3 in this case, resulting in the following automaton:



If you need further convincing that this is equivalent to our original DFA, try constructing a bi-simulation.

Note that for NFAs, the same result does *not* hold. Consider the two automata below:



These two automata both recognise the same language – aa^* – and are minimal for this language, but they are not isomorphic.

3.10 Limitations of regular languages

So far we have assumed implicitly that all languages of interest are regular, that is, that they can be described using regular expressions or finite automata. This is not really true, but the reason regular languages are so important is that they have nice descriptions, and that they suffice for many purposes.

Something a finite automaton cannot do is to count—or at least, it cannot count beyond a bound defined *a priori*. Such an automaton has a finite number of states, and its only memory is given by those states. Hence it cannot remember information that cannot be encoded into that many states.

If we want an automaton that decides whether or not a word consists of at least three letters then this automaton needs at least four states: One to start in, and three to remember how many letters it has already seen. Similarly, an automaton that is to decide whether a word contains at least 55 *a*s must have at least 56 states.

However, if we try to construct an automaton that decides whether a word contains precisely as many 0s as 1s, then we cannot do this: Clearly the automaton must have a different state for every number of 0s it has already seen, but that would require it to have infinitely many states, which is not allowed.

Similarly, how would one construct a pattern for the language

$$\mathcal{L} = \{0^n 1^n \mid n \in \mathbb{N}\}?$$

We can certainly cope with the language $\{(01)^n \mid n \in \mathbb{N}\}$ by using the pattern $(01)^*$, but that is because once we have seen 0, and the subsequent 1, we may forget about it again (or return to the start state of our automaton). In order to describe \mathcal{L} , on the other hand, we really have to remember how many 0s there are before we find the first 1. But there could be any number of 0s, so our automaton would have to have arbitrarily many states—or, if we built a pattern $\epsilon|01|0011|000111|\dots$ it would be infinite. This is not allowed.

There is a result that allows one to conclude that a particular language is not regular, the **Pumping Lemma**. Looking at that is beyond the scope of this course, but if you really want to understand this issue you should look this result up. There are automata that do have a memory device which allows them to cope with such languages, but we do not treat them here.

Exercise 65. *Regular Languages*

Which of the following languages are regular? If you answer yes, prove it. If it is no, try to say why.

- (a) The language of all words that contain at least as many *a*s as *b*s.
- (b) The language of all strings of brackets that are balanced (so scanning the word from left to right we cannot have more closing brackets at any time than we have seen opening brackets, and for the whole word the number of closing brackets is equal to the number of opening brackets).
- (c) The language of all strings over the alphabet $\{0, 1\}$ which, when viewed as a binary number, are divisible by 3.

Exercise 66. *Regular Languages*

Repeat Exercise 65 for the following languages.

- (a) The language of all strings over the alphabet $\{0, 1\}$ which, when viewed as a binary number, are prime numbers.
- (b) The language of all non-empty strings over the alphabet $\{0, 1, \dots, 9\}$ such that the last symbol is one that occurred earlier in the string.
- (c) The language of all strings over $\{0, 1\}$ where the number of 0 is different from the number of 1s.
- (d) The language of all palindromes over $\{0, 1\}$, that is, all strings that read forwards as backwards or, to put it differently, the strings which are equal to their reversal.
- (e) The language of all strings over $\{0, 1\}$ such that there are two 0s in the string which are separated by a number of characters that is a non-zero multiple of 4.
- (f) The language of all odd-length palindromes over $\{a, b, c\}$.

- (g) The language of all words over the alphabet $\{a, b\}$ that have the same number of occurrences of the substring ab as that of the substring ba . The word bab does belong to that language, since it has one occurrence of the substring ab ($[ab]a$) and one of the substring ba ($a[ba]$). *Does your answer change when the alphabet is changed to $\{a, b, c\}$?
- (h) The set of all words of the form $s1^n$, where s is a word over $\{a, b, c\}$, and n is the number of letters in s .

3.11 Summary

- We can use *finite automata* to describe the same languages as those given by regular expressions, that is, the regular languages.
- Automata can be *deterministic* (DFA) or *non-deterministic* (NFA). Every DFA can be viewed as an example of an NFA, and for every NFA we can construct a DFA that recognizes precisely the same words.
- Given a DFA we can construct a regular expression such that the language defined by this expression is precisely the language recognized by the automaton we started with.
- Given a regular expression we can construct an *NFA with ϵ -moves* such that the words accepted by the automaton are precisely those that match the regular expression. We can then eliminate the ϵ -moves from the NFA, and use the algorithm mentioned above to turn that into a DFA that recognizes the same language.
- As a result, what we can express using regular expressions, or DFAs, or NFAs, are precisely the same languages.
- We can form the union, intersection, complement, reversal, concatenation and Kleene star operation on regular languages and we will then get another regular languages.
- Not all languages are regular. In particular those languages for which an automaton would have to count up to a boundless number are not regular.

Chapter 4

Describing more complicated languages

At the end of Section 3 it is demonstrated that regular expressions are not sufficient to describe all the languages we are interested in. In particular we cannot even generate the language of all balanced brackets—clearly that is something we have to do when trying to cope with programming languages. How is a compiler to break up a Java program if it cannot cope with checking whether the braces `{, }` line up properly?

In this section we introduce a way of talking about languages that is more general. However, it is also more complicated, so whenever possible it is a good idea to stick with regular languages.

4.1 Generating words

Instead of thinking of patterns that may be matched by strings it is our aim here to *generate* all the words in our language following some well-defined rules. The way of reading the rules is to see them as ways of *rewriting the string generated so far*, by using one of the rules given to change one symbol into a string of others. These rules are given by what we call a *grammar*.

Consider the following example:

$$\begin{array}{ll}
 B \rightarrow 0 & S \rightarrow B \\
 B \rightarrow 1 & S \rightarrow (S) \\
 B \rightarrow 2 & S \rightarrow S + S \\
 B \rightarrow 3 & S \rightarrow S - S \\
 B \rightarrow 4 & S \rightarrow S \times S \\
 B \rightarrow 5 & S \rightarrow S / S \\
 B \rightarrow 6 & \\
 B \rightarrow 7 & \\
 B \rightarrow 8 & \\
 B \rightarrow 9 &
 \end{array}$$

We have to know where to start, and we always use S as the only symbol we may create from nothing. Once we have S , we can use any of the rules on the right hand side to replace it by a more complicated string, for example

$$\begin{aligned}
 S &\Rightarrow S \times S \Rightarrow S \times (S) \Rightarrow S \times (S - S) \Rightarrow B \times (S - S) \Rightarrow 5 \times (S - S) \\
 &\Rightarrow 5 \times (B - S) \Rightarrow 5 \times (B - B) \Rightarrow 5 \times (3 - B) \Rightarrow 5 \times (3 - 4).
 \end{aligned}$$

We call this a **derivation** from the grammar.

Exercise 67. Derivations

Write out derivations for $3/4 + 5$, and 9. Can you write a derivation of 12? Write out a different derivation for $3/4 + 5$. How many of these can you see?

In other words, a grammar is a set of rules that allows us to build certain expressions in a purely mechanical manner. We could program these rules into a computer and ask it to begin to generate strings from it, but unless we were careful regarding what exactly we told it it might build one very long string until it ran out of memory.

Note that the only words we are really interested in are those which consist entirely of the digits, and the arithmetic operators—they no longer contain either S or B . We merely need these two symbols to help us build the string in question: They can be seen as placeholders for longer expressions that haven't yet been built. We call the letters we want our final strings to be built of *terminal symbols* (because they are the only ones there at the end). Once we have introduced a terminal symbol it cannot be replaced by another. The other symbols are called *non-terminal symbols* (because they are not terminal). They are also sometimes called *variables*. Think of them as auxiliary symbols, or placeholders: Once they have done their job they disappear. The rules for the grammar only apply to non-terminal symbols—there can never be a rule that allows us to rewrite a terminal symbol.

Exercise 68. Grammars

Try to write down a grammar with terminal symbols $\{0,1\}$ for all non-empty words that start with 0 and end with 1. Show how your grammar generates the strings 01 and 001 and convince yourself that it does not allow you to generate invalid strings such as 11 or 00. *Hint: If you find that difficult at this stage then wait until you have seen the formal definition for a grammar and another example or two.*

We find the above description of a grammar a bit lengthy, and often use an abbreviated form of giving the rules that tell us how we may rewrite a string. For the example this looks as follows.

$$\begin{aligned} B &\rightarrow 0 \mid 1 \mid 2 \mid 3 \mid 4 \mid 5 \mid 6 \mid 7 \mid 8 \mid 9 \\ S &\rightarrow B \mid (S) \mid S + S \mid S - S \mid S \times S \mid S / S \end{aligned}$$

What we have done is to group together all the ways in which we may replace either of the two symbols B and S .

4.2 Context-free grammars

It is time for a formal definition.

Definition 18. A *context-free grammar* or **CFG** is given by the following:

- an alphabet Σ of terminal symbols, also called the object alphabet;
- an alphabet Ξ of non-terminal symbols, the elements of which are also referred to as auxiliary characters, placeholders or, in some books, variables, with $\Xi \cap \Sigma = \emptyset$;
- a special non-terminal symbol $S \in \Xi$ called the start symbol;
- a finite set of production rules of the form $R \rightarrow X$ where $R \in \Xi$ is a non-terminal symbol and $X \in (\Sigma \cup \Xi)^*$ is an arbitrary string of terminal and non-terminal symbols, which can be read as ' R can be replaced by X '.^a

^aThis assumes that we have a special symbol, namely \rightarrow which should not be contained in either Σ or Ξ to avoid confusion.

In the example above, the object alphabet is

$$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, +, \times, /, -, (,)\}.$$

The non-terminal symbols are B and S .

Why do we call grammars whose production rules satisfy these conditions *context-free*? It is because the rules tell us that we may replace any occurrence of a non-terminal as directed. A *context-sensitive grammar* would be able to have rules such as ‘You may replace the non-terminal A by the string (A) provided that there is an S immediately to the left of A ’, so a production rule might look like this: $SA \rightarrow S(A)$.

In general, what context-sensitivity means is that according to their production rules strings consisting of terminal and non-terminal symbols can be replaced by other such strings. The advantage of context-freeness is that we know that once a terminal character is in place, it will remain in the word even if further reduction rules are applied. Further, we don’t have to worry about having to choose between different production rules taking us in different directions to the same degree.

Assume we had rules

$$AB \rightarrow A \quad \text{and} \quad AB \rightarrow B.$$

Given AB occurring in some string we would then have to make a choice whether to keep the A or the B . For context-free grammars, if we have two distinct non-terminal symbols (such as A and B here) we can apply any one rule for A , and any one rule for B , and it does not matter which order we do that in.

If we allowed context-sensitive rules then our grammars would be much more complicated, and while it would be possible to express more complicated features it would be much harder to see, for example, which language a grammar generates. Natural languages certainly require context-sensitive grammars, but when it comes to programming we can go a very long way without having to use this idea.

The following definition says how we may generate strings from a grammar, and how a grammar defines a language.

Definition 19. A string $X \in (\Sigma \cup \Xi)^*$ is a **string generated by grammar G** if there is a sequence of strings

$$S = X_0 \Rightarrow X_1 \Rightarrow \cdots \Rightarrow X_n = X$$

such that each step $X_i \Rightarrow X_{i+1}$ is obtained by the application of one of G ’s production rules to a non-terminal occurring in X_i as follows.

Let $R \in \Xi$ occur in X_i and assume that there is a production rule $R \rightarrow Y$. Then X_{i+1} is obtained from X_i by replacing one occurrence of R in X_i by the string Y .

The **language generated by grammar G** is the set of all strings over Σ which are generated by G .^a

^aNote that such a string may not contain any non-terminal symbols.

Let us look at another example. Let us assume that we want to generate a grammar for the language of all strings that are non-empty and start with 0 over the alphabet $\{0, 1\}$. How do we go about it? Well, we know what our underlying alphabet of terminal symbols Σ is going to be. How about non-terminal ones? This is harder to decide *a priori*. We know that we need at least one of these, namely the start symbols S . Whether or not further non-terminal symbols are required

depends on what it is that we have to remember when creating a word (just as the number of states required for a finite automaton depends on what we need to remember about a word¹).

How do we make sure that our word starts with 0? Well, we have a production rule for S that says $S \rightarrow 0 \cdots$ then we certainly create 0 as the left-most symbol, and this will stay the left-most symbol. But what should the dots be?

If we try $S \rightarrow 0S$ then we do get a problem: Eventually we have to allow 1s into our string, but if we have a rule that says $S \rightarrow 1S$ then we could generate words whose first letter is 1. An obvious solution to this problem is to use a new non-terminal symbol, say A . Then we can have a rule

$$S \rightarrow 0A,$$

and all we now have to do is to allow A to create whatever letters it likes, one at a time, or to vanish into the empty string. Hence we want rules

$$A \rightarrow 0A \mid 1A \mid \epsilon.$$

Exercise 69. *Single Non Terminal*

Can you create a CFG for the same language using just one non-terminal symbol? *Hint: If you find this tough now then come back to it after you have had more practice.*

Exercise 70. *Grammars*

This exercise asks you to create CFGs for various languages over the alphabet $\{0, 1\}$.

- (a) All words that contain two consecutive 1s. Now give a derivation for the word 0111. Can you create a grammar for this language with just one non-terminal symbol?
- (b) All words that contain at least two 1s. Now give a derivation for 01010.
- (c) All words whose length is precisely 3.

Which ones of those are easy to do, which ones harder? Look at other languages from Exercises 11, 12, 13, 14 and 15. Which ones of those, do you think, would be particularly hard?

Exercise 71. *Grammars*

Repeat Exercise 70 for the following.

- (a) All words whose length is at most 3.
- (b) All words whose length is at least 3.
- (c) All words whose length is at least 2 and whose last but one symbol is 0. Now give a derivation for 0101.
- (d) All words for which every letter at an even position is 0. Give a derivation of 101, and of ϵ .
- (e) All words consisting of an arbitrary number of 0s followed by the same number of 1s. Try to explain how to generate a word of the form $0^n 1^n$ where $n \in \mathbb{N}$.
- (f) All words that contain an even number of 0s. Make sure that your language can generate strings such as 001010 and 0001001000.

¹However, in a finite state machine we only get to scan the word from left to right; when creating a derivation we may be working on several places at the same time.

- (g) All words that do not contain the string 01.

Parts of Exercise 71 already indicate that we can express languages using CFGs that we cannot express using finite automata or regular expressions.

We give another example. Assume we want to generate the set of all words of the language

$$\{s1^n \mid s \in \{a, b, c\}^*, n \in \mathbb{N} \text{ and } n \text{ is the number of letters in } s\}.$$

This cannot be a regular language since it would require an automaton to remember how many letters it has seen, and that number might rise above any bound (compare Exercise 65 and 66). Our grammar has to be such that every letter is matched by 1, so we have to create them simultaneously. If we use $\Sigma = \{a, b, c, 1\}$ and $\Xi = \{S\}$ the following production rules will do just that.

$$S \rightarrow aS1 \mid bS1 \mid cS1 \mid \epsilon.$$

Here is another exercise that asks you to find grammars for other non-regular languages (compare this with Exercise 65 and 66).

Exercise 72. Grammars

Give CFGs for the following languages.

- (a) The language of all strings of brackets that are balanced (so we cannot have more closing brackets at any time than we have seen opening brackets, and for the whole word the number of closing brackets is equal to the number of opening brackets).
- (b) The language of all palindromes over $\{a, b, c, d\}$.

Exercise 73. Grammars

Repeat Exercise 72 for the following.

- (a) The language of all odd-length palindromes over $\{a, b, c\}$ whose middle letter is a .
- (b) The language

$$\{a^i b^j c^k \mid i, j, k \in \mathbb{N}, k = i + j\}.$$

Exercise 74. Grammars

Repeat Exercise 72 for the following.

- (a) The language of all even-length palindromes over $\{a, b, c\}$.
- (b) The language of words over $\{a, b, c\}$ defined as

$$\{s^i c^j \mid s = ab, i, j \in \mathbb{N}, j \geq i\}.$$

We have a name for the languages that can be described by a context-free grammar.

Definition 20. A language \mathcal{L} over Σ is **context-free** if there is a context-free grammar whose alphabet of terminal symbols is Σ such that \mathcal{L} is generated by the grammar

Exercise 75. *Counting*

Consider the following grammar.

$$\Sigma = \{a, b\}, \Xi = \{S\}$$

$$S \rightarrow SaSaSbS \mid SaSbSaS \mid SbSaSaS \mid \epsilon$$

Argue that this CFG generates the language of all words that contain precisely twice as many *a*s as *b*s.

Exercise 76. *Counting*

Design a CFG for the language consisting of all words that contain precisely as many *a*s as *b*s. Try to argue that you really can generate all such words using your production rules. *Hint: You could analyse the language from the previous exercise. Alternatively, try to come up with a way that keeps track of the number of *a*s or *b*s generated so far.*

4.3 Regular languages are context-free

We have already seen that we can capture non-regular languages using context-free grammars. It may seem fairly clear that CFGs are more powerful than regular expressions, but how do we know that we can express *every* regular language using a CFG?

This turns out not to be so very difficult. Maybe surprisingly it is easier here to start with a DFA than with a regular expression.

Assume that we have a regular language \mathcal{L} over some alphabet Σ . By Theorem 3.5 we can find a DFA, say $(Q, q_\bullet, F, \delta)$, that recognizes this language.

We now define a CFG from the automaton as follows:

- We use Σ as the alphabet of terminal symbols.
- We use $\Xi = Q$ as the alphabet of non-terminal symbols.
- We use $S = q_\bullet$ as the start symbol.
- For every transition


$$q \xrightarrow{x} q'$$

in the automaton (that is, for every pair $(q, x) \in Q \times \Sigma$ and $q' = \delta(q, x)$) we introduce a production rule

$$q \rightarrow xq'.$$

- For every accepting state $q \in F$ we add a production rule $q \rightarrow \epsilon$.

We recall the following automaton from Section 3 where we have changed the names of the states.



We can now use non-terminal symbols S and A as well as terminal symbols 0 and 1 for a CFG with the following production rules:

$$\begin{aligned} S &\rightarrow \epsilon \mid 1S \mid 0A \\ A &\rightarrow 1A \mid 0S \end{aligned}$$

This grammar will generate the language of those words consisting of 0s and 1s which contain an even number of 0s.

Exercise 77. Grammars

Use this construction of a context-free grammar to produce a grammar for the regular language defined by the pattern $(ab|a)^*$. *Hint you should already have a DFA for this pattern from Exercise 51.*

Exercise 78. Grammars

Use this construction of a context-free grammar to produce a grammar for the regular language defined by the pattern $((00)^*11|01)^*$. *Hint you should already have a DFA for this pattern from Exercise 52.*

We note that the production rules we have used here are very limited: for an arbitrary non-terminal symbol $R \in \Xi$ they are of the form $R \rightarrow xR'$, where $x \in \Sigma$ and $R' \in \Xi$, or $R \rightarrow x$, where² $x \in \Sigma$, or $R \rightarrow \epsilon$. We call grammars where all production rules are of this shape **right-linear**. Such grammars are particularly simple, and in fact every language generated by a right-linear grammar is regular.

We can do even more for regular languages using CFGs. This is our first truly applied example of a context-free grammar.

There is a context-free grammar for the language of regular expressions over some alphabet Σ .

The underlying alphabet of terminal symbols³ we require here is

$$\Sigma \cup \{\epsilon, \emptyset, |, *\}.$$

The alphabet of non-terminal symbols can be $\{S\}$. We use the following production rules:

- $S \rightarrow \emptyset$.
- $S \rightarrow \epsilon$.
- $S \rightarrow x$, for all $x \in \Sigma$.
- $S \rightarrow (SS)$ for concatenation.
- $S \rightarrow (S|S)$ for alternative.
- $S \rightarrow (S^*)$ for the Kleene star.


4.4 Parsing and ambiguity

When a compiler deals with a piece of code it has to *parse* it. In other words, it has to break it down into its constituent parts in a way that makes sense. Similarly, if you were to write a program that could take (possibly quite long) arithmetic expressions as input and evaluate it, it would have to break down the expression to decide in which order to carry out the various operations.

When we give a derivation we automatically provide one way of breaking down the given string. Parsing is an attempt of carrying out the *opposite* process, namely to take a string and find out how to assemble it from simpler parts. Instead of finding a derivation it is often more meaningful to create a **parse tree** which gives a better overview of how the various bits fit together.

²We did not need that rule above, but we need it to be part of this definition.

³It would not be a good idea to call this alphabet Σ this time.



We can read off the string generated from the leaves of the tree⁴. Here it is $5 \times (3 - 4)$, and we can also see which production rules to use to get this result. If we compare this with a derivation like that given on page 77 we can see that in general, there will be several derivations for every parse tree. This is because the parse tree does not specify in which order the rewriting of non-terminal symbols should be carried out.

If we have $S \times S$, should we first replace the left S by B , or the right one by (S) ? The answer is that it doesn't really matter, since we can do this in either order, and the strings we can get to from there are just the same. What the parse tree does then is to remember the *important* parts of a derivation, namely which rules to apply to which non-terminals, while not caring about the order some of these rules are applied.

Any non-trivial word in a context-free language typically has many different derivations. However, ideally we would like every word to only have one parse tree for reasons explained below.

Consider the string $5 \times 3 - 4$. This has the following two parse trees.



Now if a computer (or a calculator) is to decide how to calculate this number it has to break down the string. The parse trees above say that this can be done in two ways, which we might write as $(5 \times 3) - 4$ and $5 \times (3 - 4)$. The first is calculated as $15 - 4 = 11$, while the second gives $5 \times (-1) = -5$, so they really do give rise to quite different things.

Of course we have generally adopted the convention that it is the first expression that we mean when we write $5 \times 3 - 4$, because the convention says that we should multiply before we add. But if we look at our grammar then there is nothing that specifies that a computer should treat this expression in this way.

If there are two different parse trees for the same word we say that it is **ambiguously generated** and we say that the grammar is **ambiguous**. When designing grammars for languages that require that the generated words are evaluated in some way then one should always aim to give a *unambiguous* grammar. However, it is not always possible to do this, and if that is the case then we call the language in question **inherently ambiguous**.

Definition 21. A context-free grammar is *ambiguous* if there is a word w with two (or

⁴Note that ϵ may appear as a leaf, in which case we ignore it from reading off the word in question.

more) *distinct parse trees* for w .

We can change the grammar given above in such a way that every word it generates can be parsed unambiguously as follows.

$$\begin{aligned} B &\rightarrow 0 \mid 1 \mid 2 \mid 3 \mid 4 \mid 5 \mid 6 \mid 7 \mid 8 \mid 9 \\ S &\rightarrow B \mid (S + S) \mid (S - S) \mid (S \times S) \mid (S/S) \end{aligned}$$

Now when we derive a word we automatically say in which order the various operations should be carried out because there are brackets doing this for us.⁵ Note that a word still has a number of different *derivations*, it is the number of *parse trees* that should be equal to one.

The first thing a compiler has to do when given a program is to recognize various key words, which are often called *tokens*. Examples for tokens for the Java programming languages are words such as `if`, `else`, `public`, `class`, and many others, but also the names of variables and methods. This is referred to as a *lexical analysis*, and this is typically carried out using regular expressions. The Unix command `lex` generates such analysers.

Once this is done the compiler has to parse the program—in other words, it produces a parse tree for the program. If that cannot be done then the compiler will raise a syntax error. There are some standard tools to help with this kind of thing: The Unix `YACC` command takes a context-free grammar and produces a parser for that grammar. Clearly for a program it is particularly important that it can be uniquely parsed so that there is precisely one way for the computer to carry out the constructions contained in it. Because programmers do not like to put in a lot of brackets typically a compiler will have a list of which operations should be carried out first.

We look at another example of ambiguity. Consider the following grammar.

$$\begin{aligned} S &\rightarrow aT \mid Ta \\ T &\rightarrow aT \mid bT \mid \epsilon \end{aligned}$$

The word aa has two parse trees for this grammar:



Hence we know that this grammar is also ambiguous. How can we turn it into an unambiguous one? First of all we have to find out which language is described by the grammar. In this case this is not so difficult: It is the set of all words over the alphabet $\{a, b\}$ which start with a or end with a (or both). The ambiguity arises from the first rule: The a that is being demanded may either be the first, or the last, letter of the word, and if the word starts and ends with a then there will be two ways of generating it.

What we have to do to give an unambiguous grammar for the same language is to pick the first or last symbol to have priority, and the other one only being relevant if the chosen one has not been matched. A suitable grammar is as follows:

$$\begin{aligned} S &\rightarrow aT \mid bU \\ T &\rightarrow aT \mid bT \mid \epsilon \\ U &\rightarrow aU \mid bU \mid a \end{aligned}$$

⁵The usual convention that tells us that multiplication should be carried out before addition is there so that we do not have to write so many brackets for an expression while it can still be evaluated in a way that leads to a unique result.

Now we generate any word from left to right, ensuring the grammar is unambiguous. We also remember whether the first symbol generated is a , in which case the remainder of the word can be formed without restrictions, or b , in which case the last symbol has to be a . It turns out that we do not need three non-terminal symbols to do this: Can you find a grammar that uses just two and is still unambiguous?

Note that in order to demonstrate that a grammar is *ambiguous*, we need only produce *one* word for which there are multiple parse trees. To demonstrate that a grammar is *unambiguous* we need to show that no such word can exist. This will usually involve some reasoning about the grammar, for example using the observation above that every word is generated from left to right and thus can only have one possible parse tree.

Exercise 79. Ambiguity

For the grammar over the alphabet $\Sigma = \{a, b\}$ given below do the following: Show that the grammar is ambiguous by finding a word that has two parse trees, and give both the parse trees. Now try to determine the language generated by the grammar. Finally give an unambiguous grammar for the same language.

Let $\Xi = \{S, T\}$ with production rules $S \rightarrow TaT$ and $T \rightarrow aT \mid bT \mid \epsilon$.

Exercise 80. Ambiguity

Repeat Exercise 79 for the grammar with $\Xi = \{S\}$ and $S \rightarrow aS \mid aSbS \mid \epsilon$.

Exercise 81. Ambiguity

Repeat Exercise 79 for the grammar with an underlying alphabet of $\{a, b, c\}$, two non-terminal symbols S and T , the start symbol S and production rules:

$$\begin{aligned} S &\rightarrow TabT \\ T &\rightarrow aT \mid bT \mid cT \mid \epsilon \end{aligned}$$

4.5 A programming language

Defining a programming language is a huge task: First we have to describe how programs can be formed (we have to define the **syntax** of the language), and then we have to describe how a computer should act when given a particular program (we have to give a **semantics**).

The syntax of a language tells us how to write programs that are considered well-formed (so if it is a compiled language, for example, the compiler will not give any syntax errors). However, this says nothing about what any program actually *means*. When you see a fragment of a program you probably have some kind of idea what at least some of the instructions within it are supposed to do, but that is because program designers stick to certain conventions (so that $a + b$ will typically mean adding the content of the variable a to that of the variable b), but a computer doesn't know anything about these preconceptions. Therefore in order to fully define a programming language one also has to say what a program actually means. There is a third year course unit that concerns itself with different ways of doing so.

Most real programming languages such as Java, C, or ML, or PHP, typically have very large grammars and we do not have the time here to study such a big example.

We look at a very small programming language called **While**. At first sight it looks as if there can't possibly anything interesting to be computable with it: It only knows about boolean data types

(such as `true`, T , here `tt`, and `false`, F , here `ff`), and natural numbers, say from 0 to some number we call `maxint`. We also need variables which can take on the value of any of our natural numbers. Let's assume for the moment we only have three variables, x , y , and z . The language has three kinds of entities:

- arithmetic expressions A
- boolean expressions B and
- statements S .

For each of these we have a production rule in our grammar.

$$\begin{aligned} A &\rightarrow 0 \mid 1 \mid \cdots \mid \text{maxint} \mid x \mid y \mid z \mid A + A \mid A \times A \mid A - A \\ B &\rightarrow \text{tt} \mid \text{ff} \mid A = A \mid A \leq A \mid B \wedge B \mid \neg B \\ S &\rightarrow x := A \mid y := A \mid z := A \mid \text{skip} \mid S; S \mid \text{if } B \text{ then } S \text{ else } S \mid \text{while } B \text{ do } S \end{aligned}$$

This may surprise you, but if we add more variables to our language⁶ then it can calculate precisely what, say, `Java` can calculate for natural numbers and booleans.

We can now look at programs in this language. Assume that the value held in the variable y is 10.

$$x := 1; \text{while } \neg y = 0 \text{ do } x := 2 \times x; y := y - 1$$

This program is intended to calculate 2^y and store it in x , assuming y holds a natural number.

Exercise 82. *A Programming Language*

We look at this example in more detail.

- Give a derivation for the program above.
- Is this grammar unambiguous? If you think it isn't, then give an example and show how this grammar might be made unambiguous, otherwise explain why you think it is.
- Give a parse tree for the above program. Can you see how that tells you something about what computation is supposed to be carried out?

4.6 The Backus-Naur form

In the literature (in particular that devoted to programming languages) context-free grammars are often described in a slightly different form to the one used here so far. This is known as the *Backus-Naur form* (sometimes also described as *Backus Normal Form* but that isn't a very accurate name). There is a simple translation process.

- Non-terminal symbols don't have to be capitals and are distinguished from other symbols by being in *italics*, or they are replaced by a descriptive term included in angled brackets \langle and \rangle .
- The rewriting arrow \rightarrow is replaced by the symbol $::=$.

We repeat the grammar for the `While` language from above to illustrate what this looks like.

Here we assume that a ranges over arithmetic expressions, **AExp**, b ranges over boolean expressions, **BExp**, and S ranges over statements, **Stat**.

⁶When people define grammars for programming languages they typically use abbreviations that allow them to stipulate 'any variable' without listing them all explicitly, but we don't want to introduce more notation at this stage.

$$\begin{aligned}
a &::= 0 \mid 1 \mid \dots \mid \mathbf{maxint} \mid x \mid y \mid z \mid a + a \mid a \times a \mid a - a \\
b &::= \mathbf{tt} \mid \mathbf{ff} \mid a = a \mid a \leq a \mid b \wedge b \mid \neg b \\
S &::= x := a \mid y := a \mid z := a \mid \mathbf{skip} \mid S; S \mid \mathbf{if } b \mathbf{ then } S \mathbf{ else } S \mid \mathbf{while } b \mathbf{ do } S
\end{aligned}$$

In the next example we assume that instead we are using $\langle \text{aexp} \rangle$ to range over **AExp**, $\langle \text{bexp} \rangle$ over **BExp**, and $\langle \text{stat} \rangle$ over **Stat** respectively.

$$\begin{aligned}
\langle \text{aexp} \rangle &::= 0 \mid 1 \mid \dots \mid \mathbf{maxint} \mid x \mid y \mid z \mid \\
&\quad \langle \text{aexp} \rangle + \langle \text{aexp} \rangle \mid \langle \text{aexp} \rangle \times \langle \text{aexp} \rangle \mid \langle \text{aexp} \rangle - \langle \text{aexp} \rangle \\
\langle \text{bexp} \rangle &::= \mathbf{tt} \mid \mathbf{ff} \mid \langle \text{aexp} \rangle = \langle \text{aexp} \rangle \mid \langle \text{aexp} \rangle \leq \langle \text{aexp} \rangle \mid \\
&\quad \langle \text{bexp} \rangle \wedge \langle \text{bexp} \rangle \mid \neg \langle \text{bexp} \rangle \\
\langle \text{stat} \rangle &::= x := \langle \text{aexp} \rangle \mid y := \langle \text{aexp} \rangle \mid z := \langle \text{aexp} \rangle \mid \mathbf{skip} \mid \langle \text{stat} \rangle; \langle \text{stat} \rangle \mid \\
&\quad \mathbf{if } \langle \text{bexp} \rangle \mathbf{ then } \langle \text{stat} \rangle \mathbf{ else } \langle \text{stat} \rangle \mid \mathbf{while } \langle \text{bexp} \rangle \mathbf{ do } \langle \text{stat} \rangle
\end{aligned}$$

4.7 Properties of context-free languages

There is a description for building new regular languages from old ones in Section 3. We can use almost any way we like to do so; set-theoretic ones such as unions, intersections and complements, as well as ones using concatenation or the Kleene star; even reversing all the words work. Context-free languages are rather more fiddly to deal with: Not all of these operations work for them.

Concatenation. This does work. We do the following. Assume that we have two grammars with terminal symbols taken from the alphabet Σ , and non-terminal symbols Ξ_1 respective Ξ_2 . We now take every symbol in Ξ_1 and put a subscript 1 onto it, and similarly for every symbol in Ξ_2 , and the subscript 2. We have now forced those two alphabets to be disjoint, so when we form their union the number of symbols in the union is the sum of the symbols in Ξ_1 and Ξ_2 . We add a new start symbol S to the set of non-terminal symbols. We now take all the production rules from the first grammar, and put subscripts of 1 onto each non-terminal symbol occurring in it, and do the same for the second grammar with the subscript 2. We add one new production rule $S \rightarrow S_1 S_2$.

Exercise 83. *Minimum Length*

Use your solution to Exercise 70 (c) to produce a grammar for the language of all words whose length is at least 6, using this procedure.

Kleene star. This looks as if it should be more complicated than concatenation, but it is not. We merely add two production rules to the grammar (if they aren't already there) for our given language, namely $S \rightarrow SS$, and $S \rightarrow \epsilon$. If we then wish to generate a word that is the n -fold concatenation of words in our language, where $n \in \mathbb{N}^+$, we start by applying the rule

$$S \rightarrow SS$$

$(n - 1)$ times, giving us n copies of S . For each one of these we can then generate the required word. If we wish to generate the empty word we can do this by applying the second new rule, $S \rightarrow \epsilon$.

Exercise 84. *Grammars*

Use your solution to Exercise 70 (c) to produce a grammar for the language of all words whose length is divisible by 3, using this procedure.

Reversal. This is quite easy. Leave the two alphabets of terminal and non-terminal symbols as they are. Now take each production rule, and replace the string on the right by its reverse. So if there is a production rule of the form $R \rightarrow 00R1$, replace it by the rule $R \rightarrow 1R00$.

Exercise 85. *Grammar Reversal*

Use this procedure to turn your solution for Exercise 68 into a grammar that generates all words over $\{0, 1\}$ that start with 1 and end with 0.

Unions. This does work.

Exercise 86. *Grammar Union*

Show that the union of two context-free languages over some alphabet is context-free.

Intersections. Somewhat surprisingly, this is not an operation that works. The intersection of a context-free language with a regular one is context-free, but examining this issue in more detail goes beyond the scope for this course.

Complements. This also does not work. Just because we have a way of generating all the words in a particular set does not mean we can do the same for all the words *not* belonging to this set.

4.8 Limitations of context-free languages

Given the idea of context-sensitive grammars mentioned at the start of this section it should not come as a surprise that there are languages which are not context-free. We do not develop the technical material here that allows us to prove this statement. In Section 3 there is an explanation why finite state automata can only do a very limited amount of counting, and context-free grammars can certainly do more than that. However, there are limitations to this as well. The language

$$\{a^n b^n c^n \mid n \in \mathbb{N}\}$$

over $\{a, b, c\}$ is not context-free.

There is a formal result, called *the Pumping Lemma for context-free languages* that describes some of the possible properties that indicate that a language fails to be context-free.

There are automata which can be used to describe the context-free languages which are known as *pushdown-automata*. These do have a memory device that takes the form of a stack. That makes them fairly complicated to cope with, and that is why they are beyond the scope of this course. Moreover, for these automata, deterministic and non-deterministic versions are not equivalent, making them even more complex to handle.

4.9 Summary

- A *context-free grammar* or *CFG* is given by two alphabets, one of *terminal*, and one of *non-terminal* symbols and by production rules that tell us how to replace non-terminals by strings consisting of both, terminal and non-terminal symbols.
- We generate words for a CFG by starting with the non-terminal start symbol S . In the word generated so far we then pick one non-terminal symbol and replace it by a string according to one of the production rules for that symbol.
- A CFG generates a language made up of all the words consisting entirely of terminal symbols that can be generated using it. The languages that can be generated in this way are known as the *context-free languages*.
- Every regular language is context-free, but there are some context-free languages that are not regular.

- Rather than a derivation for a word we can give a *parse tree* which is sufficient to tell us how to carry out calculations, for example. Programming languages are defined using grammars, and a compiler parses a program to check whether it is correctly formed.
- We can apply the union, concatenation, Kleene star and reversal operations to context-free languages and that gives us another context-free language.
- There are languages that are not context-free.

Appendix A

A Mathematical Approach

The account given so far is informal as far as possible. To be completely rigorous the language of mathematics has to be used to define the various notions we have used. For the interested reader we here give a glimpse of how this is done. It is possible to pass this course, and pass it with a good mark, without knowing the content of this chapter, but those who want to get the whole picture, and get a really good mark, should work through this part as well.

A.1 The basic concepts

We begin by giving mathematical definitions for the various notions from Section 2.1.

Definition 22. An **alphabet** Σ is a finite set of characters (primitive indecomposable symbols) called the **letters** of Σ .

Definition 23. Let Σ be an alphabet. A **word over** Σ is a finite string

$$x_1x_2 \cdots x_n$$

where $n \in \mathbb{N}$ and the x_i are letters from Σ . If $n = 0$ we write ϵ for the resulting string.

Exercise 87. Words

When we are given an alphabet we can look at various kinds of words we can build from it.

- (a) Form all the three letter words over the alphabet $\{a, b, 1\}$.
- (b) Form all the words over the alphabet $\{0\}$.
- (c) Form all the words over the alphabet $\{\}$.

We can alternatively use a recursive definition for a word. This is particularly useful if we want to define operations on words recursively. It also fits well with how some data structures are defined.

Alternative definition. A **word** over the alphabet Σ is obtained from the following rules:

- The empty word ϵ is a word.

- If s is a word and $x \in \Sigma$ then sx is a word.

Now that we have the basic concept we can define operations for words, such as the length of a word. Note how this definition follows the recursive definition of a word.

Definition 24. *The length $|s|$ of a word s over some alphabet is defined as follows:*

$$|s| = \begin{cases} 0 & s = \epsilon \\ |s'| + 1 & s = s'x \end{cases}$$

Exercise 88. *Non Recursive Length*

Try to come up with a non-recursive definition of the length function. *Hint: Look at the original definition of word, or at the definition of concatenation to get an idea.* Use the recursive definition of a word to argue that your definition agrees with the original.

We have a binary operation on words, namely concatenation.

Definition 25. *Given an alphabet Σ , concatenation is an operation from pairs of words to words, all over Σ , which, for word s and t over Σ , we write as $s \cdot t$. It is defined as follows:*

$$x_1 \dots x_m \cdot y_1 \dots y_n = x_1 \dots x_m y_1 \dots y_n.$$

Exercise 89. *Operators*

Recall the definition of an associative or commutative operation from COMP11120.

- Argue that the concatenation operation is associative.
- Show that the concatenation operation is not commutative.

Exercise 90. *Recursive Concatenation*

Use recursion to give an alternative definition of the concatenation operation. (You may find this difficult—try to give it a go anyway.)

Exercise 91. *Length*

Show that $|s \cdot t| = |s| + |t|$. *Hint: You may well find it easier to use the non-recursive definition for everything.*

Exercise 92. *Recursion*

Practice your understanding of recursion by doing the following.

- Using the recursive definition of a word give a recursive definition of the following operation. It takes a word, and returns the word where every letter is repeated twice. So ab turns into $aabb$, and aba into $aabbaa$, and aa to $aaaa$.
- Now do the same thing for the operation that takes a word and returns the reverse of the word, so abc becomes cba .

In order to describe words of arbitrary length concisely we have adopted notation such as a^3 for aaa . This works in precisely the same way as it does for powers of numbers: By a^3 we mean the word that results from applying the concatenation operation to three copies of a to obtain aaa , just as in arithmetic we use 2^3 to indicate that multiplication should be applied to three copies of 2 to obtain $2 \times 2 \times 2$. So all we do by writing a^n is to find a shortcut that tells us how many copies of a we require (namely n many) without having to write them all out.

Because we know both these operations to be associative we do not have to use brackets here: $(2 \cdot 2) \cdot 2$ is the same as $2 \cdot (2 \cdot 2)$ and therefore the notation $2 \cdot 2 \cdot 2$ is unambiguous, just as is the case for $2 \times 2 \times 2$.

What should a^1 mean? Well, this is simple, it is merely the word consisting of one copy of a , that is a . The question of what a^0 might mean is somewhat trickier: What is a word consisting of 0 copies of the letter a ? A useful convention in mathematics is to use this to mean the unit for the underlying operation. Hence in arithmetic $2^0 = 1$. The unit for concatenation is the empty word and so we think of a^0 as a way of referring to the empty word ϵ .

This way we obtain useful rules such as $a^m \cdot a^n = a^{n+m}$. Similarly we have $(a^m)^n = a^{nm}$, just as we have for exponentiation in arithmetic. However, note that $(ab)^n$ consists of n copies of ab concatenated with each other, rather than $a^n b^n$, as we would in arithmetic.¹

Exercise 93. *Words*

Write out in full the words 0^5 , $0^3 1^3$, $(010)^2$, $(01)^3 0$, 1^0 .

Languages are merely collections of words.

Definition 26. *Let Σ be an alphabet. A language over Σ is a set of words over Σ .*

As mentioned in Chapter 2 using this definition we automatically obtain set-theoretic operations: We can form the unions, intersections, complements and differences of languages in precisely the same way as we do this for other sets. Hence expressions such as $\mathcal{L}_1 \cap \mathcal{L}_2$, $\mathcal{L}_1 \cup \mathcal{L}_2$ and $\mathcal{L}_1 \setminus \mathcal{L}_2$ are immediately meaningful.

Exercise 94. *Language Operators*

Let Σ be the alphabet $\{0, 1, 2\}$ and let

$$\mathcal{L}_1 = \{s \mid s \text{ is a word consisting of 0s and 1s only}\},$$

$$\mathcal{L}_2 = \{s \mid s \text{ is a word beginning with 0 and ending with 2}\}.$$

Calculate the following languages: $\mathcal{L}_1 \cap \mathcal{L}_2$, $\mathcal{L}_1 \cup \mathcal{L}_2$ and the complement of \mathcal{L}_1 in the language of all words over Σ .

Note that we have notation to find a more compact description of the language

$$\{\epsilon, 1, 11, 111, 1111, \dots\}$$

of all words over the alphabet $\{1\}$ as

$$\{1^n \mid n \in \mathbb{N}\}.$$

Exercise 95. *Sets*

Write down the following languages using set-theoretic notation:

- (a) All the words consisting of the letters a and b which contain precisely two a s and three b s.

¹The reason this rule doesn't hold is that the concatenation operation isn't commutative (see Exercise 89), and so we can't swap over the a s and b s to change the order in which they appear.

- (b) All the words consisting of the letter 1 that have even length.
- (c) All the words consisting of an arbitrary number of a s followed by at least one, but possibly more, b s.
- (d) All the non-empty words consisting of a and b occurring alternately, beginning with an a and ending with a b .
- (e) All the non-empty words consisting of a and b occurring alternately, beginning with an a and ending with an a .
- (f) All the words consisting of an arbitrary number of 0s followed by the same number of 1s.
- (g) All the words over some alphabet Σ .

Note that if we have two alphabets $\Sigma' \subseteq \Sigma$ then every word over Σ' can be viewed as a word over the alphabet Σ . Alternatively, we may restrict a language \mathcal{L} over Σ to all those words that only use letters from Σ' by forming

$$\{s \in \mathcal{L} \mid s \text{ only uses letters from } \Sigma'\}.$$

The definitions in Chapter 2 for forming new languages from existing ones are rigorous as given. You may want to improve your understanding of them by carrying out the following exercises.

Exercise 96.

- (a) Calculate $\{a, b\}^2$ and $\{ab, ba\}^2$.
- (b) Find the shortest set-theoretic expression you can find for the set of all words over the alphabet $\{0, 1, 2\}$ which consist of precisely three letters.
- (c) Calculate $\{00, 11\}^*$.
- (d) If $\Sigma_1 \subseteq \Sigma_2$ are two alphabets, argue that $\Sigma_1^* \subseteq \Sigma_2^*$.

A.2 Regular expressions

The formal definition of a pattern, and that of a word matching a pattern is given in Chapter 2. There are exercises in that chapter that encourage you to explore the recursive nature of the two definitions. The following should not be considered a formal exercise. Just think about it.

Exercise 97. *Non Matching Patterns*

What would it take to demonstrate that a given word does *not* match a given pattern?

We have already seen that different patterns may describe the same language. In fact, we can come up with a number of rules of rewriting regular expressions in such a way that the language they define stays the same. Here are some examples, where p , p_1 and p_2 are arbitrary patterns over some alphabet Σ .

- $\mathcal{L}(\epsilon p) = \mathcal{L}(p)$.
- $\mathcal{L}(p_1 | p_2) = \mathcal{L}(p_2 | p_1)$.
- $\mathcal{L}(p^*) = \mathcal{L}(\epsilon | p p^*)$.
- $\mathcal{L}((p_1 | p_2) p) = \mathcal{L}(p_1 p | p_2 p)$.

Exercise 98. Rules

Come up with three further such rules.

Exercise 99. Equalities

Argue that the equalities used in the recursive definition of a language matching a pattern on page 18 have to hold according to Definition 3.

Exercise 100.

For Exercise 10 find set-theoretic notation for the sets you described in English.


Exercise 101. Proposition 2.1

Prove Proposition 2.1. *Hint: Take a closer look at the recursive definition of the language defined by a pattern above.*

A.3 Finite state automata

Matching up the formal definition of an automaton with its image is fairly straightforward.

As an example consider the picture on page 23 of an automaton that accepts precisely those words over $\{0, 1\}$ that contain an even number of 0s.



We have two states, which in the picture are labelled E and O , so the set of all states is $\{E, O\}$. The initial state is E , and there is only one accepting state which is also E .

The transition function δ for this example is given by the following table:

input	output
$(E, 1)$	E
$(E, 0)$	O
$(O, 1)$	O
$(O, 0)$	E


If we follow a word such as 001 through the automaton we can now do this easily: We start in state E and see a 0, so we next need to go to $\delta(E, 0)$ which is O . The next letter is 0 again, and from the new state O we need to go to $\delta(O, 0) = E$. The last letter is 1, so from the current state E we need to go to $\delta(E, 1) = E$. Since E is an accepting state we accept the word 001.

Exercise 102. Describing Pictures

Describe the final picture from Section 3.1 in this way.

When moving to non-deterministic automata we no longer have a transition function, but a transition relation. Therefore we cannot use a table as above to write it out.

As an example, take the non-deterministic automaton from page 29.



It has states $\{0, 1, 2, 3\}$, start state 0 and accepting states $\{3\}$. The transition relation δ is defined as follows.

δ relates		to			
state	letter	0	1	2	3
0	0		✓		
0	1			✓	
1	0		✓		
1	1		✓		✓
2	0			✓	
2	1			✓	
3	0			✓	
3	1			✓	

Think of the ticks as confirming that from the given state in the left hand column there is an edge labelled with the letter in the second column to the state in the top row. We can see that δ here is a relation rather than a function because for the input state 1, letter 1, we find two ticks in the corresponding row.

Exercise 103. *Finding Patterns*

Draw a non-deterministic automaton for the language described in Exercise 11 (c). Then describe it in the same way as the above example.

Exercise 104. *Drawing Automata*

Draw the automata with states $\{0, 1, 2, 3\}$, start state 0, accepting states $\{0, 2\}$ and the following transitions.

(a)

δ relates		to			
state	letter	0	1	2	3
0	<i>a</i>		✓		
0	<i>b</i>			✓	
1	<i>a</i>				✓
1	<i>b</i>			✓	
2	<i>a</i>			✓	
2	<i>b</i>				✓
3	<i>a</i>				✓
3	<i>b</i>				✓

(b)

δ relates		to			
state	letter	0	1	2	3
0	<i>a</i>	✓	✓		
0	<i>b</i>	✓			
1	<i>a</i>	✓			✓
1	<i>b</i>	✓		✓	
2	<i>a</i>		✓		
2	<i>b</i>				
3	<i>a</i>	✓			
3	<i>b</i>		✓	✓	

Can you say anything about the resulting automata?

Exercise 105. *Algorithm 1*

Using Algorithm 1 on Page 36 we know how to turn a non- deterministic automaton into a deterministic one. Write out a definition of the transition function for the new automaton.

Exercise 106. *Algorithm 1*

Prove Theorem 3.1. All that remains to be established is that a DFA generated by an NFA using Algorithm 1 from page 36 accepts precisely the same words as the original automaton.

Pictures are good and well when it comes to describing automata, but when such a description has to be given to a computer it is best to start from the mathematical description and turn it into a data structure.

Exercise 107. *Java*

What would it take to define an automaton class in **Java**? How would a class for deterministic automata differ from one for non-deterministic ones? *Hint: You may not know enough Java yet to decide this. If this is the case then return to the question later.*

Exercise 108. *Algorithm 2*

Algorithm 2 is quite complicated. To try to understand it better I recommend the following for the general description which starts on page 45.

- Justify to yourself the expressions for $\mathcal{L}_{j \rightarrow i}^{\leq j}$ and $\mathcal{L}_{j \rightarrow i}^{\leq i}$ from the general ones for $\mathcal{L}_{j \rightarrow i}^{\leq k}$.
- Justify the equalities given at the end of the description of the general case.

Exercise 109. *Dump States*

Explain why it is safe not to draw the ‘dump’ states for the two automata when constructing an automaton for the intersection of the languages recognized by the automata as described on page 60 by forming the product of the two automata.

Exercise 110. *Reversal Recap*

Exercise 54 has two parts that are mathematical, namely a and (e). If you have delayed answering these do so now.

Let us now turn to the question of when two automata recognize the same language. We begin by giving a proper definition of what it means for two automata to be isomorphic.

Definition 27. An automaton $(Q, q_\bullet, F, \delta)$ over the alphabet Σ is isomorphic to an automaton $(P, p_\bullet, E, \gamma)$ over the same alphabet if and only if there exists a function $f: Q \longrightarrow P$ such that

- f is bijective;
- $f(q_\bullet) = p_\bullet$;
- q is in F if and only if $f(q)$ is in E for all $q \in Q$ and all $p \in P$;
- in case the automata are
 - deterministic: $f(\delta(q, x)) = \gamma(f(q), x)$ for all $q \in Q$ and all $x \in \Sigma$;
 - non-deterministic: $q \xrightarrow{x} q'$ according to δ is true if and only if $f(q) \xrightarrow{x} f(q')$ according to γ also holds.

Exercise 111. *Isomorphism*

Convince yourself that you understand this definition! Also convince yourself that two automata are isomorphic precisely when they have the ‘same’ picture (if the states remained unlabelled).

In section 3.9 we had a calculation used to turn the problem of equivalence of two automata into one of deciding whether two (different) automata recognize any words at all. You should be able to prove this, using definitions from the set-theory part of COMP11212.

Exercise 112. *Set Equalities*

For subsets S and S' of some set T show that $S = S'$ if and only if $S \cap (T - S') = \emptyset$ and $S' \cap (T - S) = \emptyset$.

The notion of bisimulation deserves additional study and, indeed, this is a concept that is very important in concurrency. Here we restrict ourselves to one more (if abstract) example.

Exercise 113. *Simulations*

- (a) Show that if two DFA accept the same language then there are simulations in both directions between them.
- (b) Find two automata that accept the same language but which do *not* have simulations in both directions between them. Hint: Will these be DFA or NFA?
- (c) Assume we are given an NFA, and the DFA that results from applying Algorithm 1 to that NFA. Show that there is always a simulation from the NFA to the corresponding DFA. Hint: You need to do Exercise 105 first. Use the ‘element’ relation between states.

Exercises 65 and 66 are somewhat mathematical in the way they demand you to think about languages. Try to do as many of their parts as you can.

A.4 The Pumping Lemma

In Section 3.10, we referred to the **Pumping Lemma**, a result that allows us to conclude that certain languages are not regular. Here, we present some more details of the Pumping Lemma as an illustration of a more formal approach to arguing and reasoning about languages. There are further details of this in Sipser, Section 1.4 and Hopcroft et.al. Section 4.1.

The Pumping Lemma gives us an important property of regular languages. For any regular language, we can find a number n such that for any word in the language with a size larger than n , there is a non-empty substring in the word that can be “pumped”, i.e. replaced by an arbitrary number of copies of itself, resulting in another word in the language.

Lemma A.1. The Pumping Lemma for Regular Languages.

For any regular language \mathcal{L} , there is a constant n such that for any word $w \in \mathcal{L}$ with $|w| \geq n$, we can break w into three strings x, y and z with $w = xyz$ and:

- (i) $y \neq \epsilon$;
- (ii) $|xy| \leq n$;
- (iii) For all $k \geq 0$, the word $xy^kz \in \mathcal{L}$.

To prove the lemma, we make use of a result known as the “pigeonhole principle”.

Proposition A.2. The Pigeonhole Principle. *If we have n things to put into m containers, where $n > m$ and n and m are finite, then we must end up with at least one container containing more than one item^a.*

^aAlthough this may seem obvious, it relies on the fact that m and n are *finite*. The pigeonhole principle is also discussed in Sipser and Hopcroft et.al.

Our proof of the lemma is as follows. Consider a regular language \mathcal{L} over Σ . We know that there must be a DFA $A = (Q, q_\bullet, F, \delta)$ such that A accepts precisely those words in \mathcal{L} . A is a finite-state automaton, so $|Q| = n$ for some $n > 0$.

Take any word $w = x_1x_2 \cdots x_m$, where $m \geq n$, $x_i \in \Sigma$ and $w \in \mathcal{L}$. We construct x, y and z .

For $i = 0, 1, \dots, n$, define the state p_i to be the state that the automaton will be in after reading the first i symbols of w . So, we have:

$$\begin{aligned} p_0 &= q_\bullet \\ p_1 &= \delta(p_0, x_1) \\ p_2 &= \delta(p_1, x_2) \\ \dots &= \dots \\ p_n &= \delta(p_{n-1}, x_n) \end{aligned}$$

We know that there are only n states in A , and we have identified $n + 1$ states here, thus by the Pigeonhole principle, it must be the case that two of these states must be the same, e.g. we can find two integers i and j such that $0 \leq i < j \leq n$ and $p_i = p_j$. These integers i and j give us the indices that we can use to break our word up. So we have:

$$\begin{aligned} x &= a_1a_2 \cdots a_i \\ y &= a_{i+1}a_{i+2} \cdots a_j \\ z &= a_{j+1}a_{j+2} \cdots a_m \end{aligned}$$

So x is the first part of the word up to the “repeated state”, y is the part of the word that the automaton consumes before arriving in the repeated state again, and z is the rest of word that is consumed before arriving in an accepting state². Figure A.1 represents the path through the automaton that is taken when it consumes the word.




Figure A.1: Path through an Automaton

We now consider the conditions listed above.

- (i) y cannot be empty as i is strictly less than j , and thus y has at least one letter in it.
- (ii) $|xy| = |a_1a_2 \cdots a_i a_{i+1} a_{i+2} \cdots a_j| = j \leq n$.
- (iii) For the third condition, we refer to Figure A.1. Consider a word xy^kz for some $k \geq 0$. If $k = 0$, the automaton consumes x and goes from start state $q_\bullet = p_0$ to p_i . As $p_i = p_j$, the automaton will then consume z and end up in an accepting state.
If $k > 0$, then the automaton will consume x and go from start state $q_\bullet = p_0$ to p_i . It will then consume the k copies of y , each time arriving back at $p_i = p_j$. Finally, it consumes z and ends up in an accepting state.

Thus our x , y and z exhibit the properties above and the lemma is proved.

Exercise 114. *Pumping Lemma for a finite language*

Demonstrate why the Pumping Lemma trivially holds for a finite language \mathcal{L} .

A.4.1 Using the Pumping Lemma

The Pumping Lemma is often used in proof by contradiction. We assume that a language is regular, and then show, using the Pumping Lemma, that a contradiction arises. Our initial assumption thus cannot hold.

Consider the language:

$$\mathcal{L} = \{0^n 1^n | n \geq 0\}$$

In Section 3.10 we made an argument that this language was not regular based on the fact that an automaton would have to “count” all of the 0s it had seen, and this cannot be done with finite number of states. We can also demonstrate this using the Pumping Lemma.

We first assume that the language *is* regular, and then show that this assumption leads to a contradiction. So, let us assume that \mathcal{L} is regular. The Pumping Lemma tells us that there is a value p such that for any word $w \in \mathcal{L}$ that is longer than p , we can split it into x , y and z such that $w = xyz$ and the conditions in the Pumping Lemma hold.

Now take the word $xyyz$. This can be seen as xy^2z , and so the Pumping Lemma says that this word should be in \mathcal{L} . We consider three different cases in turn and show that each of them arrive

²Note that there may be *other* repeated states, but we don’t care about those.

at a contradiction. We note that 1/ all words in \mathcal{L} must have an equal number of 0s and 1s and 2/ words in \mathcal{L} cannot have a 1 followed a 0.

- (i) $y = 0^n$ for some $n > 0$. The word xyz must have the same number of 0s and 1s. Thus the string $xyyz$ will have more 0s than 1s (as we have inserted a word containing just 0s and y is non-empty). It can therefore not be in \mathcal{L} and we have a contradiction.
- (ii) $y = 1^n$. Similar reasoning tells us that $xyyz$ cannot be in \mathcal{L} .
- (iii) y contains both 0s and 1s. In this case, the string $xyyz$ must have a 1 followed by a 0, and it cannot be in \mathcal{L} .

The three cases above cover all eventualities, and each one arrives at a contradiction. Thus our initial assumption – that \mathcal{L} is regular – cannot hold.

A.5 Grammars

Exercises in this section that asks you to reason mathematically are **Exercises 75** and **76**—they ask you to demonstrate something.

In Section 4.3 there is a method that takes an automaton and turns it into a context-free grammar. How can we see that the words generated by this grammar are precisely those accepted by the automaton?

If we have a word $x_1x_2 \cdots x_n$ accepted by the automaton then we get the following for the grammar.

In the automaton:

We start in state q_\bullet and go to

$$q_1 = \delta(q_\bullet, x_1).$$

Once we have reached state

$$q_i = \delta(q_{i-1}, x_i),$$

we go to

$$q_{i+1} = \delta(q_i, x_{i+1})$$

and keep going until we reach

$$q_n = \delta(q_{n-1}, x_n).$$

Now q_n is an accepting state.

For the grammar:

We have $S = q_\bullet \Rightarrow x_1q_1$.

Once we have reached the word

$$x_1x_2 \cdots x_iq_i$$

we apply a production rule to obtain

$$x_1x_2 \cdots x_ix_{i+1}q_{i+1},$$

where

$$q_{i+1} = \delta(q_i, x_{i+1}),$$

and keep going until we reach

$$x_1x_2 \cdots x_nq_n.$$

Since we know that q_n is an accepting state we may now use the production rule $q_n \rightarrow \epsilon$ to obtain $x_1x_2 \cdots x_n$.

Exercise 115. Grammars

Use a similar argument to reason that if s is a word over Σ generated by the grammar then it is also accepted by the automaton.

Exercise 116. *Right Linear Grammars*

We have already demonstrated that every regular language can be generated by a right-linear grammar. Try to show the opposite, that is, that every language generated by a right-linear grammar is regular.

Glossary

accept

A word can be accepted, or not accepted, by an automaton.

Algorithm 1

Turns an NFA into a DFA.

Algorithm 2

Turns an automaton (works for NFA and DFA, but more complicated for NFA) into a pattern describing the language recognized by the automaton.

Algorithm 3

Takes a pattern and turns it into an NFA with ϵ -transition such that the language recognized by the automaton is that defined by the pattern.

Algorithm 4

Turns an NFA with ϵ -transitions into an NFA without these.

alphabet

A set of letters.

ambiguous

A word can be ambiguously generated if we have a grammar that has two parse trees for it, and if such a word exists we call the grammar ambiguous. A language is inherently ambiguous if all the grammars generating it are ambiguous.

Backus-Naur form

A particular way of writing down a grammar popular within computer science, in particular for describing the syntax of programming languages.

CFG

Rules for generating words by using rewrite rules known as a grammar, in our case a context-free one.

concatenation

An operation on words (or letters) that takes two words and returns one word by attaching the second word to the end of the first word.

context-free

A language is context-free if there is a context-free grammar generating it. (Note that grammars can also be context-free—we do not consider other grammars on this course.)

context-free grammar

Same as CFG.

derivation

A way of demonstrating that a particular word is generated by a grammar.

deterministic finite automaton

Same as DFA.

DFA

An alternative way of defining a language, more suitable for human consumption than a regular expression.

ϵ

The empty word.

equivalence

Defined for automata. Two automata are equivalent if and only if they recognize the same language.

GNFA

Generalised Nondeterministic Finite Automaton. A special kind of automaton that can be used in deriving regular expressions from automata..

Kleene star

An operation allowing an arbitrary number of concatenations of words from a given language.

language

A set of words.

language defined by pattern

The set of all the words matching the pattern. See also the definition on page.

language generated by grammar

The set of all words consisting entirely of terminal symbols which can be built using the rules of the grammar.

letter

Building block for word. Same as symbol.

match

A word can match a pattern.

NFA

An alternative way of defining a language, often the easiest for people to devise.

NFA with ϵ -transitions

An NFA which also allows transitions labelled with ϵ which do not have to be matched against a letter.

non-deterministic finite automaton

Same as NFA.

non-terminal symbol

An auxiliary symbols used to describe a grammar—we want to create words that do not contain auxiliary symbols.

parse tree

A parse tree shows how a string can be generated from a given grammar in a more useful way than a derivation.

pattern

A particular kind of string used to define languages.

recognize

Automata recognize languages.

regular

A language is regular if it can be defined using a pattern.

regular expression

Same as pattern.

right-linear

A grammar is right-linear if its production rules are particularly simple. The definition is given on page.

simulation

A special relation between the states of two automata; its existence implies that the language accepted by the first automaton is a subset of that accepted by the second one..

string

Same as word.

string generated by grammar

A word consisting of terminal and non-terminal symbols that we can build by using the rules of the grammar.

symbol

Same as letter.

terminal symbol

An element of the alphabet over which we are trying to create words using a grammar.

unambiguous

A grammar is unambiguous if it isn't ambiguous.

word

Obtained from concatenating 0 or more letters. Same as string.